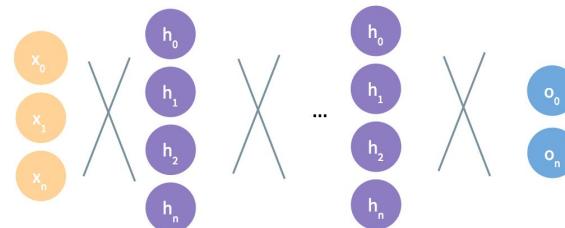


Intro to Deep Learning

Nick Locascio

6.S191



MIT 6.S191 | Intro to Deep Learning | IAP 2017



2016: year of deep learning



2016: The Year That Deep Learning Took Over the Internet

WIRED - Dec 25, 2016

The project is still in the early stages, but it hints at the widespread impact of **deep learning** over past year. In 2016, this very old but newly ...

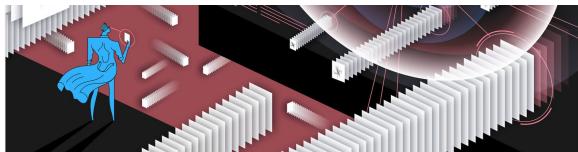


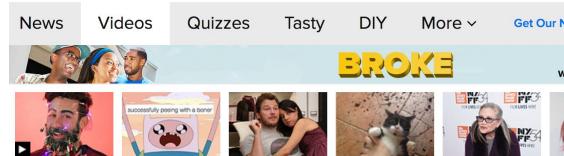
Illustration by Sandbox Studio, Chicago with Ana Kova

Deep learning takes on physics

12/06/16 | By Molly Olmstead

Can the same type of technology Facebook uses to recognize faces also recognize particles?

BuzzFeedNEWS



This Is Why A Computer Winning At Go Is Such A Big Deal

People didn't think this would happen for at least 10 years; it's a sign of how far artificial intelligence has come.

posted on Mar. 14, 2016, at 8:46 a.m.



Artificial Intelligence cybernetics machine learning technology film

Kristen Stewart co-authored a paper on style transfer and the AI community lost its mind

Posted Jan 19, 2017 by [John Mannes \(@JohnMannes\)](#)



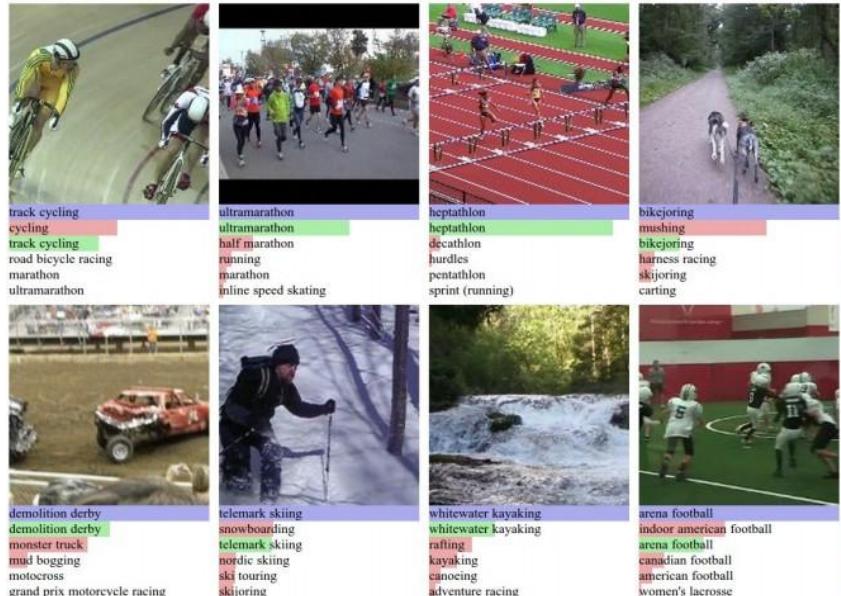
Next Story



Deep Learning Success

- Image Classification
- Machine Translation
- Speech Recognition
- Speech Synthesis
- Game Playing

... and many, many more

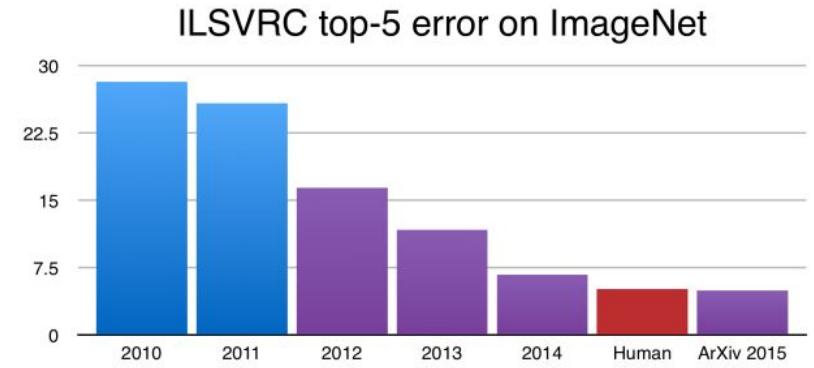


Deep Learning Success

- Image Classification

Machine Translation
Speech Recognition
Speech Synthesis
Game Playing

... and many, many more



AlexNet

Better than
humans

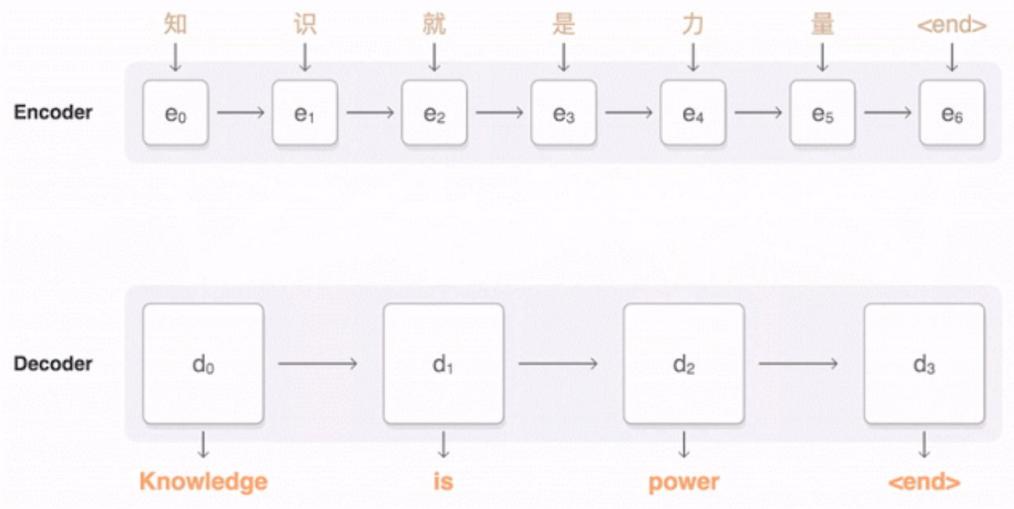
Krizhevsky, Sutskever, Hinton 2012

Deep Learning Success



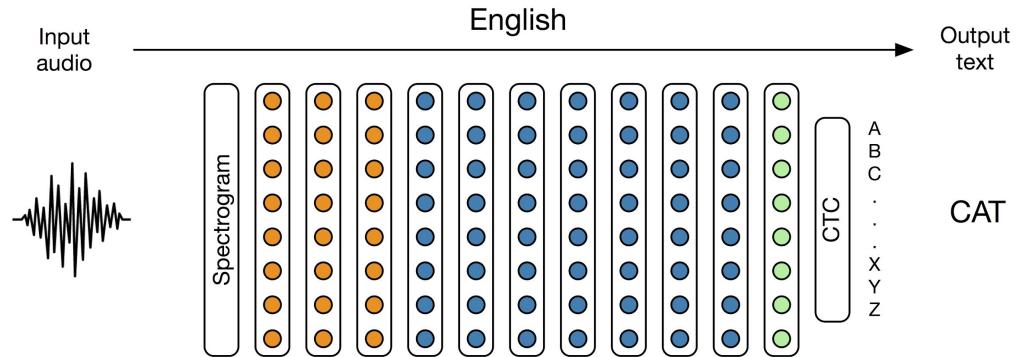
Image Classification
- Machine Translation
Speech Recognition
Speech Synthesis
Game Playing

.
... and many, many more

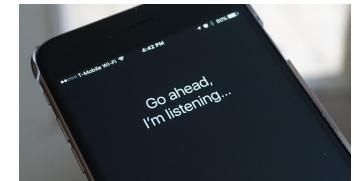


Deep Learning Success

Image Classification
Machine Translation
- **Speech Recognition**
Speech Synthesis
Game Playing



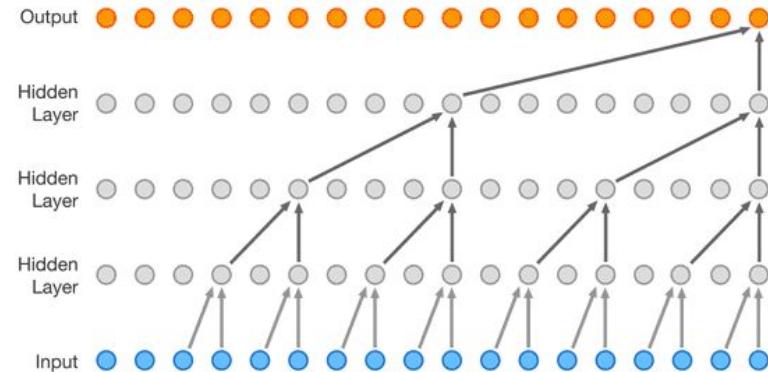
...
... and many, many more



Deep Learning Success

Image Classification
Machine Translation
Speech Recognition
- Speech Synthesis
Game Playing

... and many, many more



Deep Learning Success

Image Classification
Machine Translation
Speech Recognition
Speech Synthesis
- Game Playing



... and many, many more

Deep Learning Success

Image Classification
Machine Translation
Speech Recognition
Speech Synthesis
- Game Playing



...
... and many, many more

6.S191 Goals

1. Fundamentals
2. Practical skills
3. Up to speed on current state of the field
4. Foster an open and collaborative deep learning community within MIT

Knowledge, intuition, know-how, and community to do deep learning research and development.

Class Information

- 1 week, 5 sessions
- P/F, 3 credits
- 2 TensorFlow Tutorials
 - In-class Monday + Tuesday
- 1 Assignment: (more info in a few slides)

Typical Schedule

- 10:30am-11:15am **Lecture #1**
- 11:15am-12:00pm **Lecture #2**
- 12:00pm-12:30pm **Coffee Break**
- 12:30pm-1:30pm **Tutorial / Proposal Time**

Assignment Information

- 1 Assignment, 2 options:
 - Present a novel deep learning research idea or application
 - **OR**
 - Write a 1-page review of a deep learning paper

Option 1: Novel Proposal

- **Proposal Presentation**
 - Groups of 3 or 4
 - Present a novel deep learning research idea or application
 - 1 slide, 1 minute
 - List of example proposals on website: introtodeeplearning.com
 - Presentations on Friday
 - Submit groups by **Wednesday 5pm** to be eligible
 - Submit slide by **Thursday 9pm** to be eligible

Option 2: Paper Review

- Write a 1-page review of a deep learning paper
 - Suggested papers listed on website introtodeeplearning.com
 - We will read + grade based on clarity of writing and technical communication of main ideas.

Class Support

- **Piazza:** <https://piazza.com/class/iwmlwep2fnd5uu>
- **Course Website:** introtodeeplearning.com
- **Lecture slides:** introtodeeplearning.com/schedule
- **Email us:** introtodeeplearning-staff@mit.edu
- OH by request

Staff: Lecturers



Nick Locascio
Lead Organizer



Harini Suresh
Lead Organizer



Ishaan
Gulrajani
Co-Chair



Victoria Dean
Co-Chair



Lex Fridman
Co-Chair



Yo Shavit
Co-Chair

Staff: TA + Admin



Eduardo
DeLeon



Jackie Xu
TA



Wengong
Jin
TA



Adam Yala
TA



Harry
Bleyan
TA



Tianxiao
Shen
TA



Alex Lenail
TA



Rue Park
Marketing



Anish
Athalye



Helen Zhou
TA



Bowen
Baker
TA



Prafulla
Dhariwal
TA

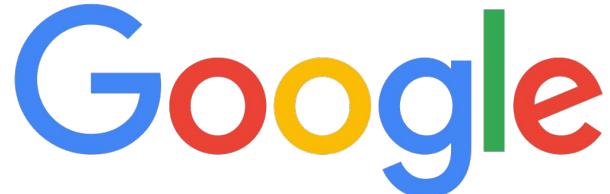


Alfredo
Yanez
TA



Hansa
Srinivasan
TA

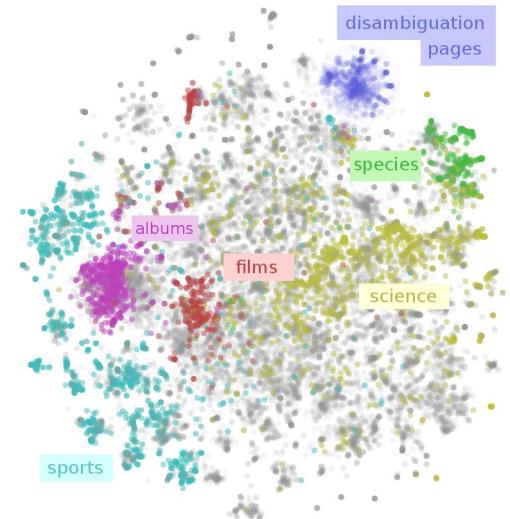
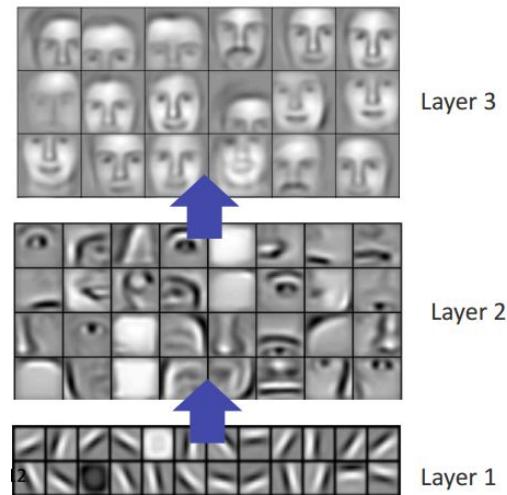
Our Fantastic Sponsors!



Why Deep Learning and why now?

Why Deep Learning?

- Hand-Engineered Features vs. Learned features



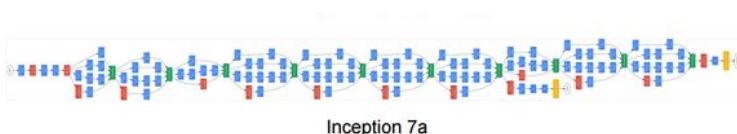
Why Now?



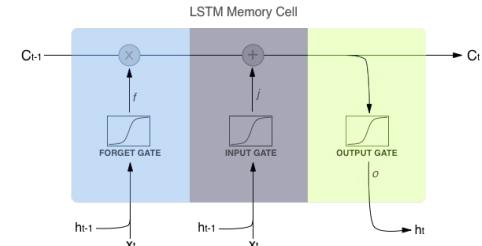
1. Large Datasets
2. GPU Hardware Advances + Price Decreases
3. Improved Techniques



WIKIPEDIA



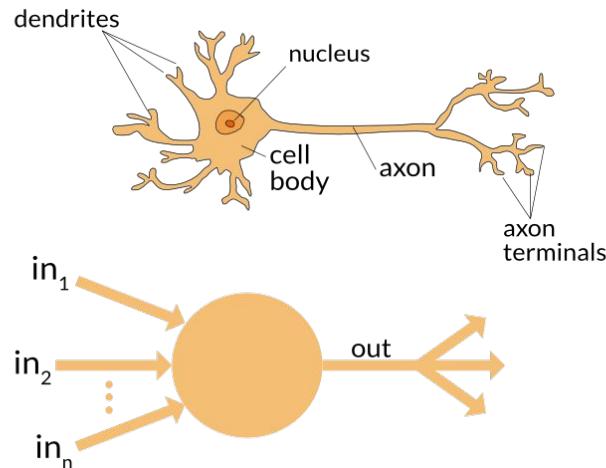
¹Going Deeper with Convolutions, [C. Szegedy et al, CVPR 2015]



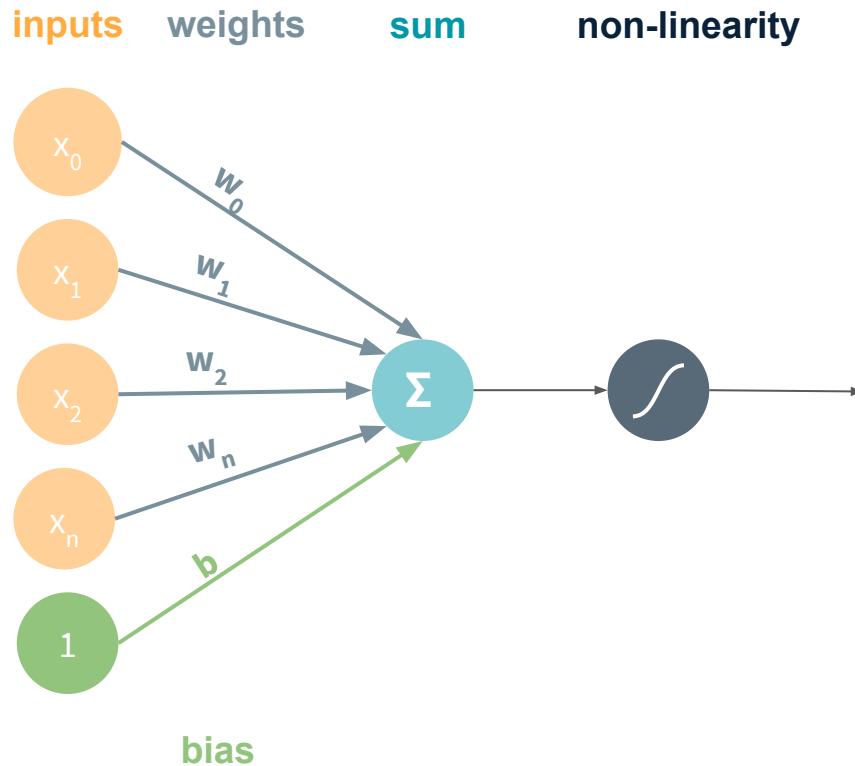
Fundamentals of Deep Learning

The Perceptron

1. Invented in 1954 by Frank Rosenblatt
2. Inspired by neurobiology

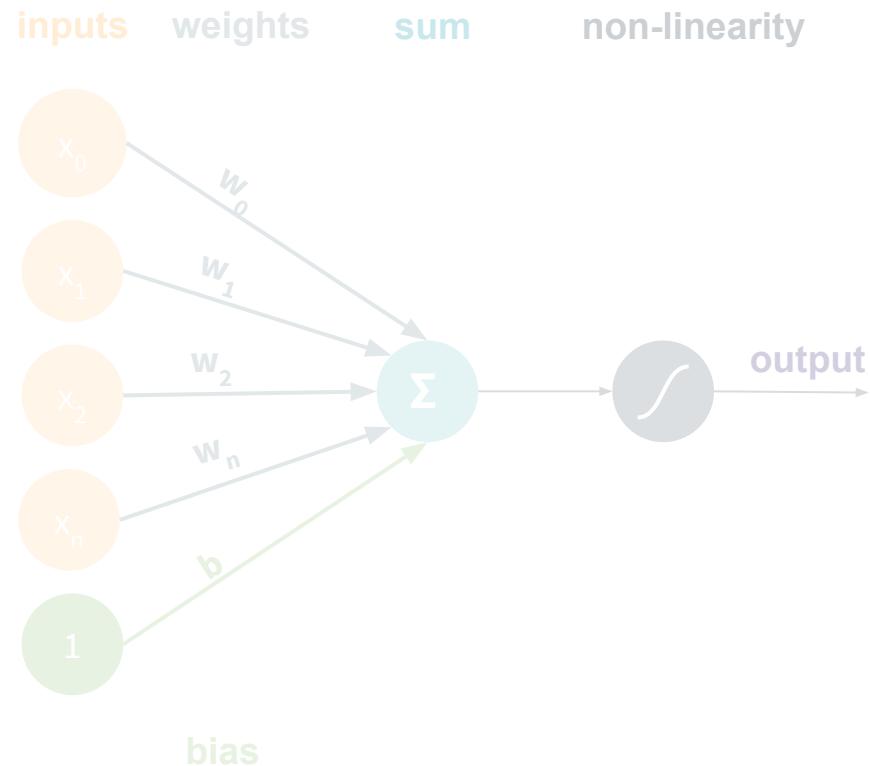


The Perceptron



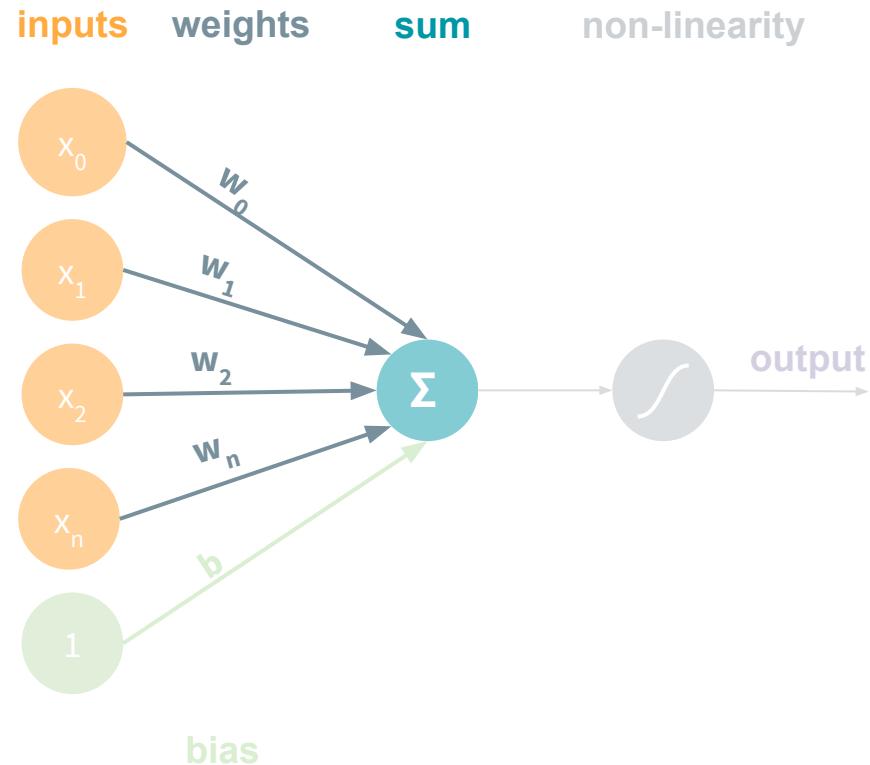
Perceptron Forward Pass

output =



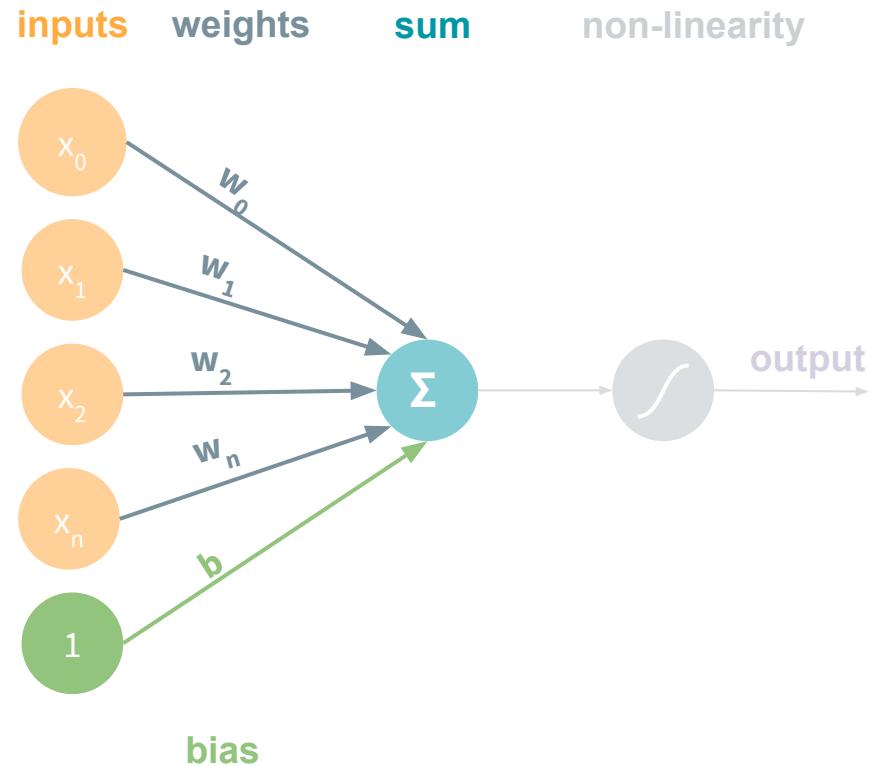
Perceptron Forward Pass

$$output = \sum_{i=0}^N x_i * w_i$$



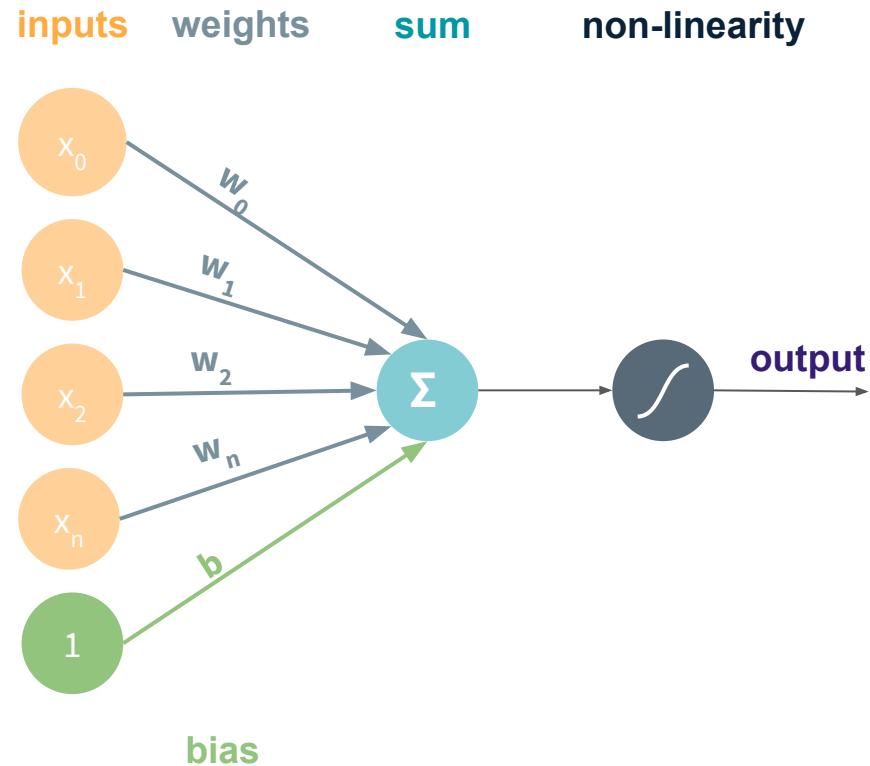
Perceptron Forward Pass

$$output = \left(\sum_{i=0}^N x_i * w_i \right) + b$$



Perceptron Forward Pass

$$output = g\left(\sum_{i=0}^N x_i * w_i\right) + b$$

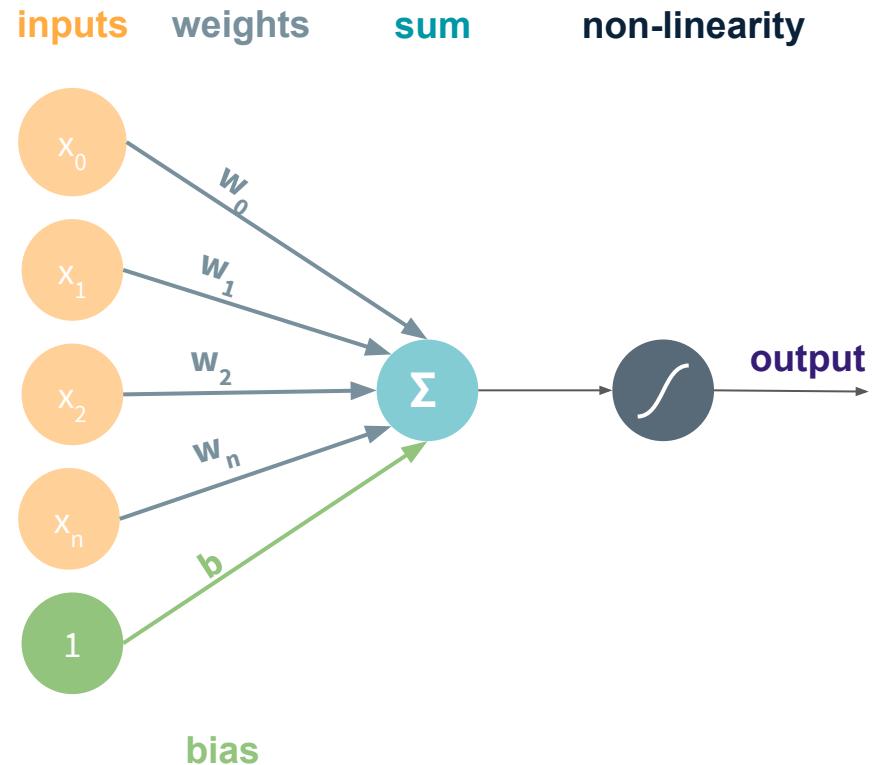


Perceptron Forward Pass

$$output = g(XW + b)$$

$$X = x_0, x_1, \dots x_n$$

$$W = w_0, w_1, \dots w_n$$



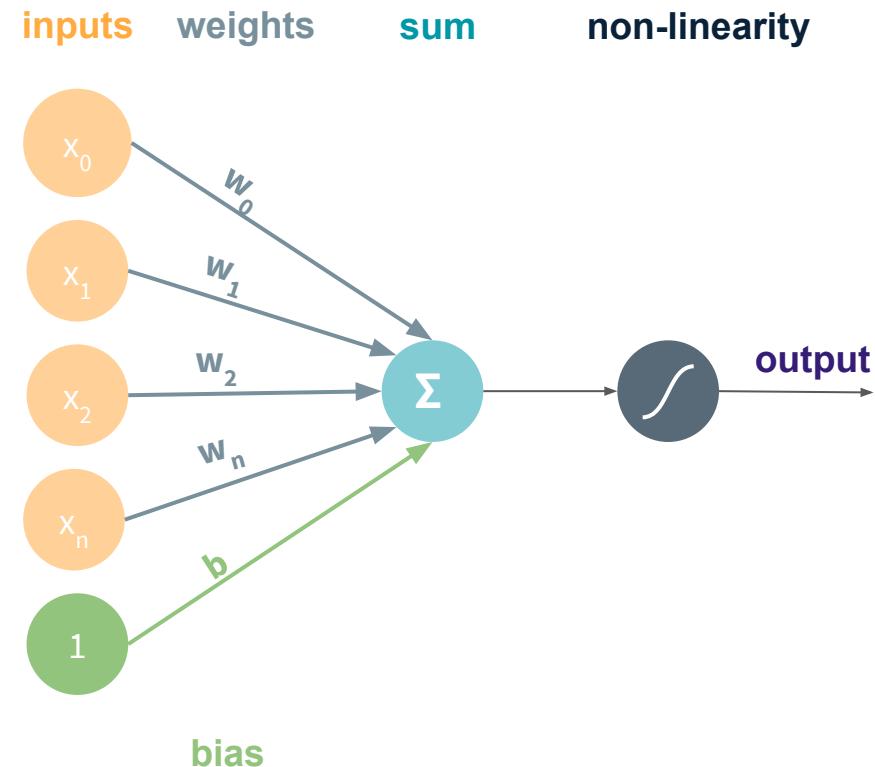
Perceptron Forward Pass

Activation Function

$$output = g(XW + b)$$

$$X = x_0, x_1, \dots x_n$$

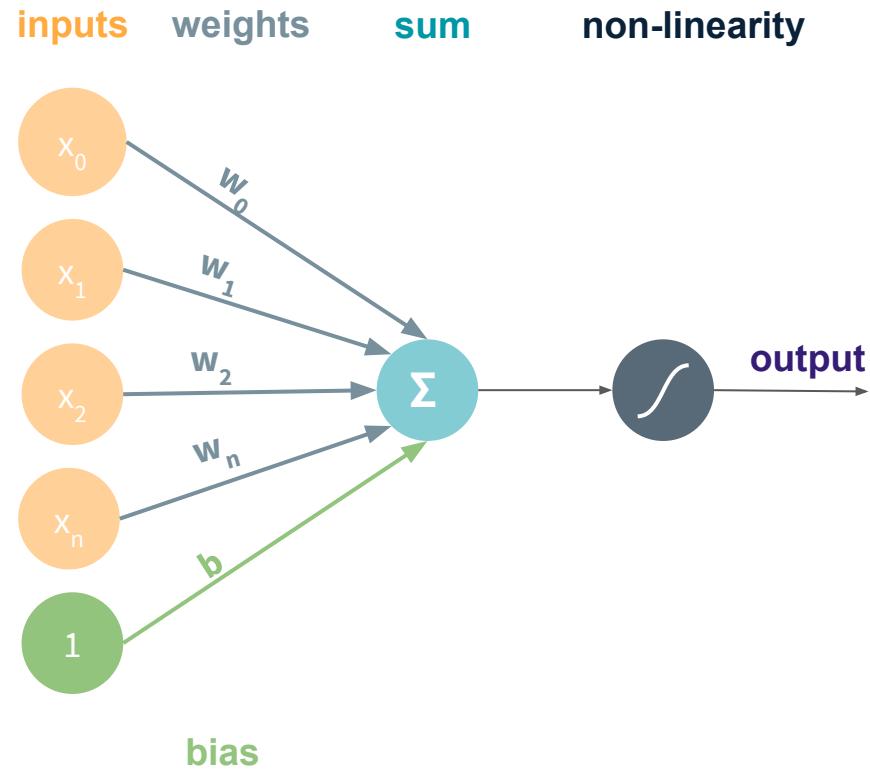
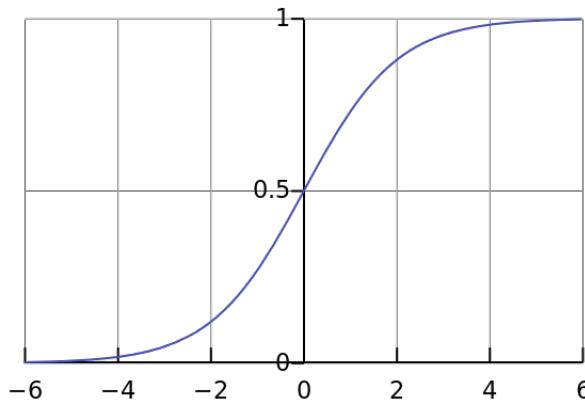
$$W = w_0, w_1, \dots w_n$$



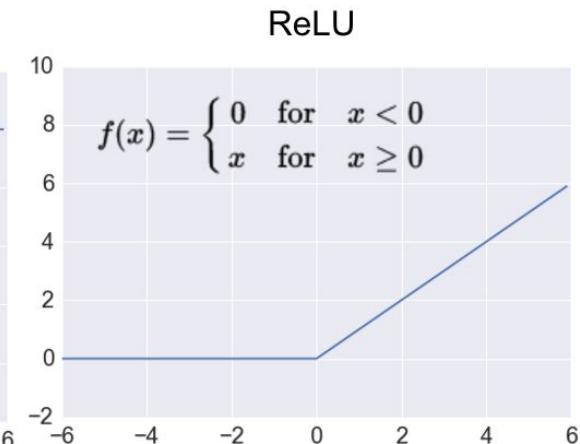
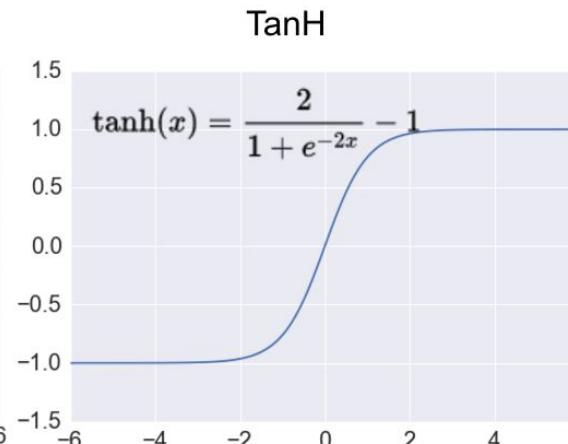
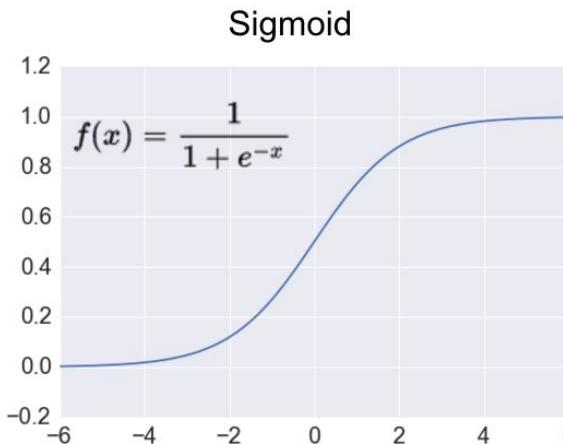
Sigmoid Activation

$$\text{output} = g(XW + b)$$

$$g(a) = \sigma(a) = \frac{1}{1 + e^{-a}}$$

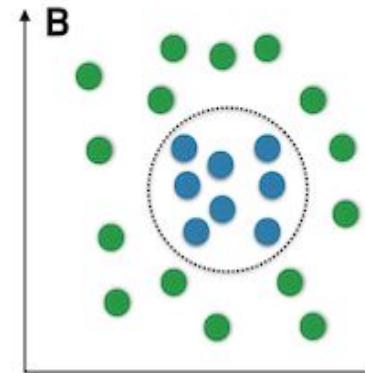
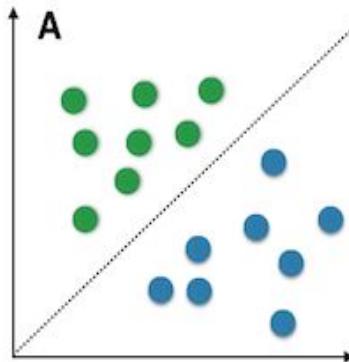


Common Activation Functions



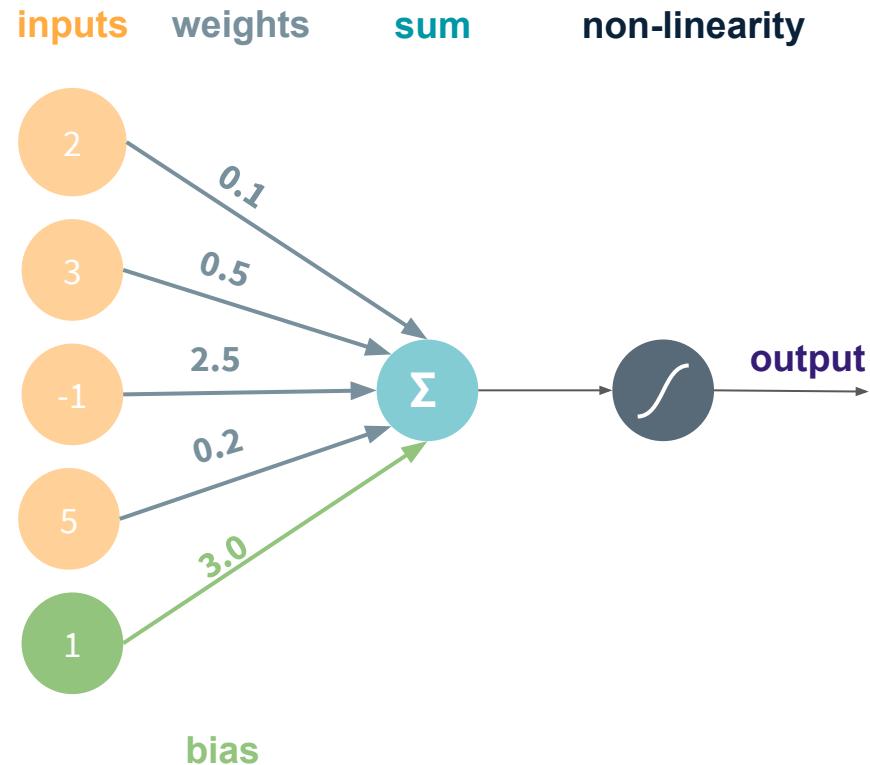
Importance of Activation Functions

- Activation functions add non-linearity to our network's function
- Most real-world problems + data are **non-linear**



Perceptron Forward Pass

$$output = g(XW + b)$$



Perceptron Forward Pass

output = $g($

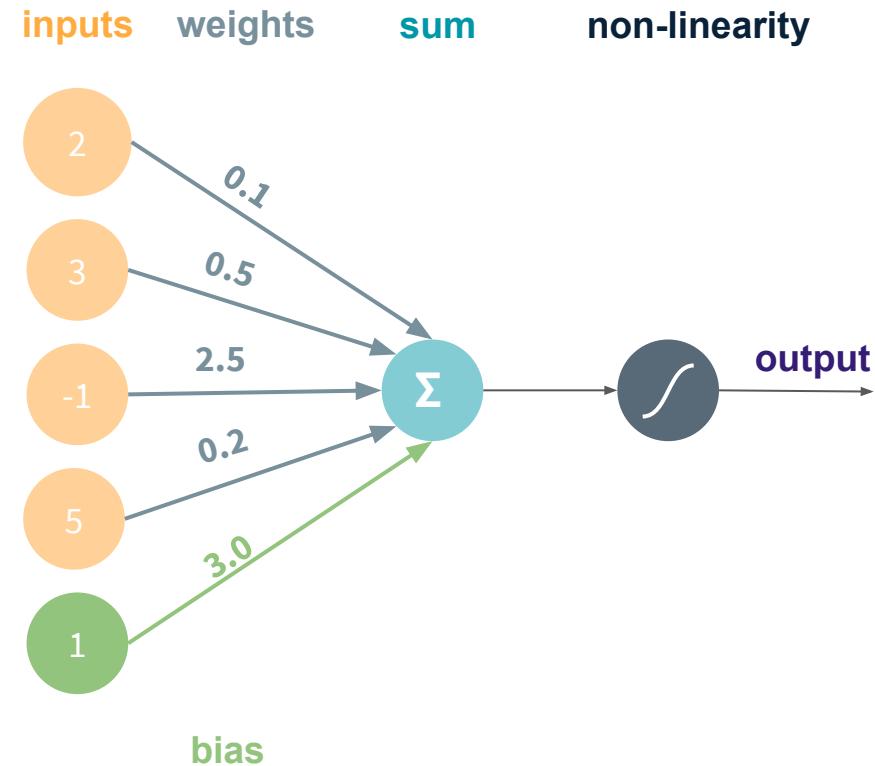
$$(2 * 0.1) +$$

$$(3 * 0.5) +$$

$$(-1 * 2.5) +$$

$$(5 * 0.2) +$$

$$(1 * 3.0)$$

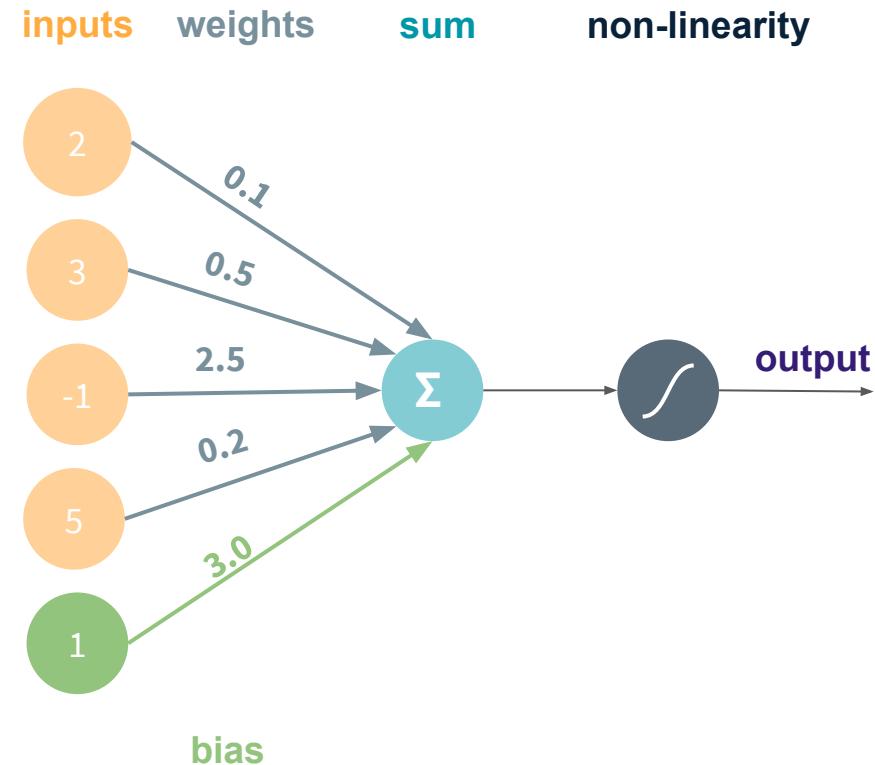


$)$

Perceptron Forward Pass

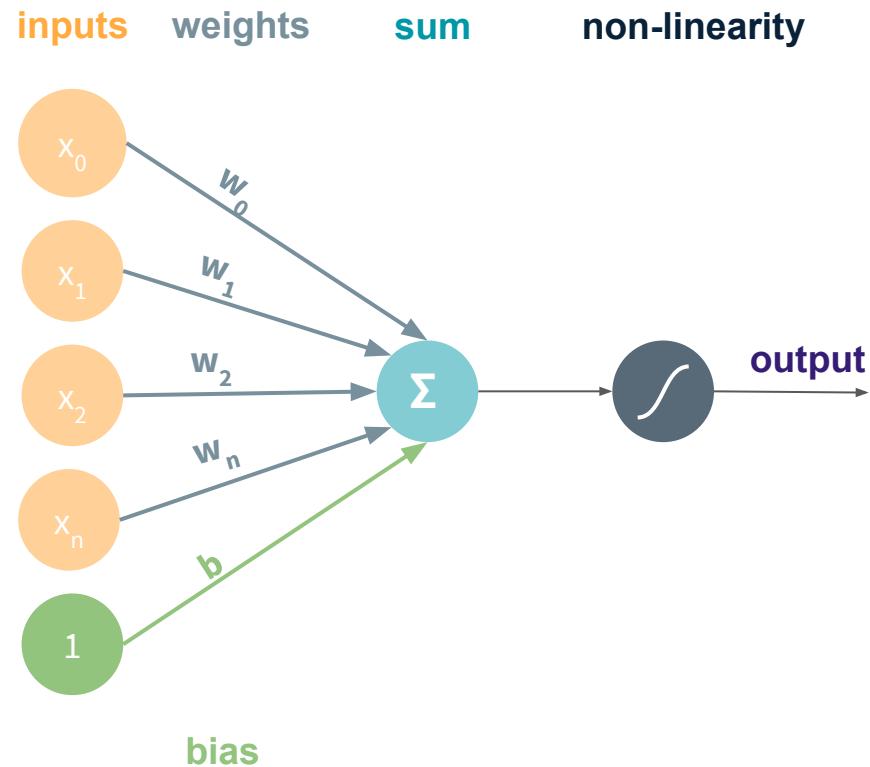
$$output = g(3.2) = \sigma(3.2)$$

$$= \frac{1}{(1 + e^{-3.2})} = 0.96$$



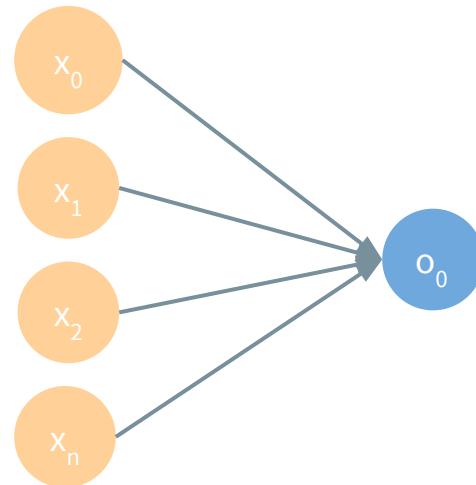
**How do we build neural networks
with perceptrons?**

Perceptron Diagram Simplified



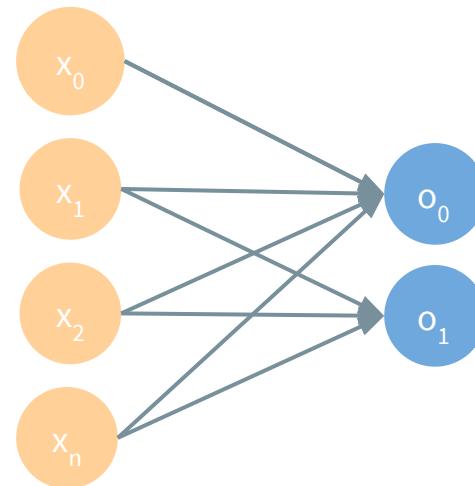
Perceptron Diagram Simplified

inputs output

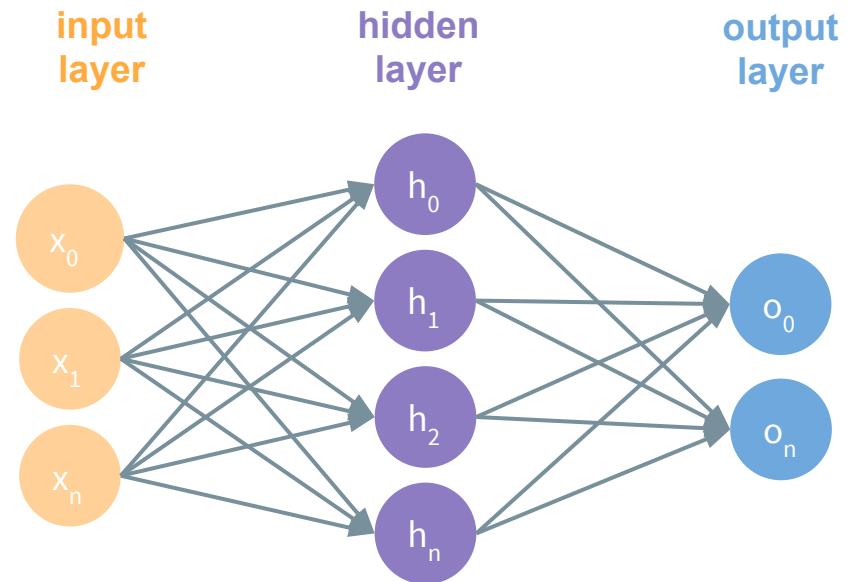


Multi-Output Perceptron

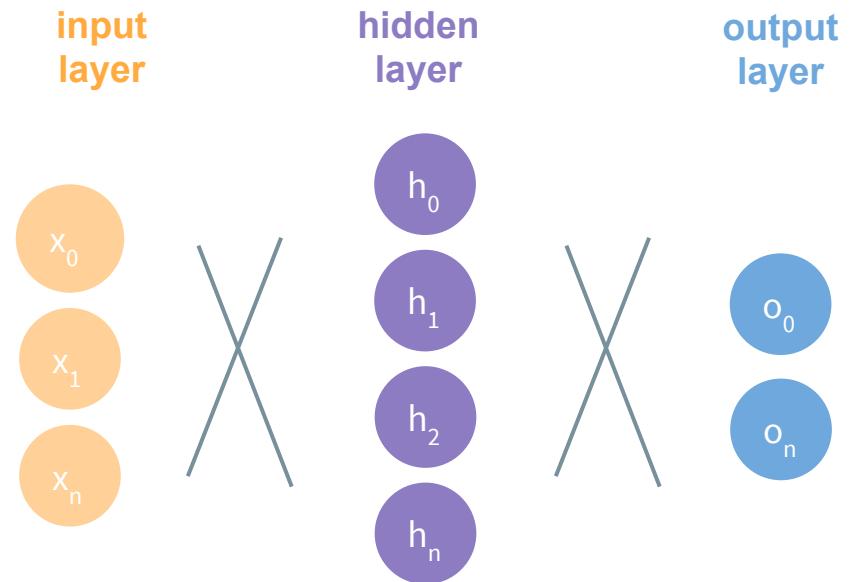
Input layer output layer



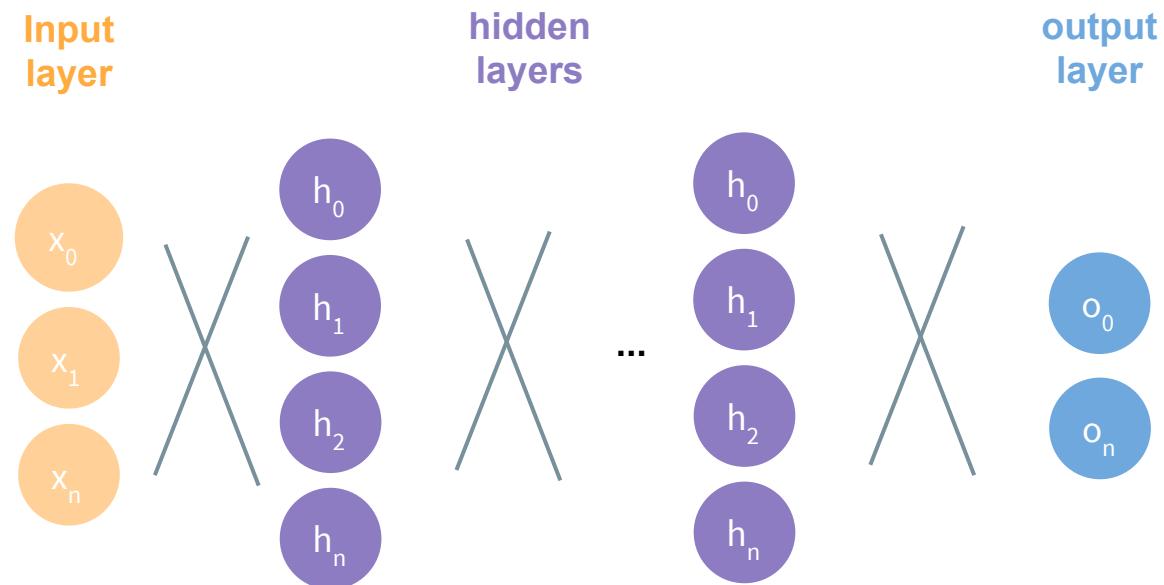
Multi-Layer Perceptron (MLP)



Multi-Layer Perceptron (MLP)



Deep Neural Network



Applying Neural Networks

Example Problem: Will my Flight be Delayed?



Example Problem: Will my Flight be Delayed?

Temperature: -20 F

Wind Speed: 45 mph



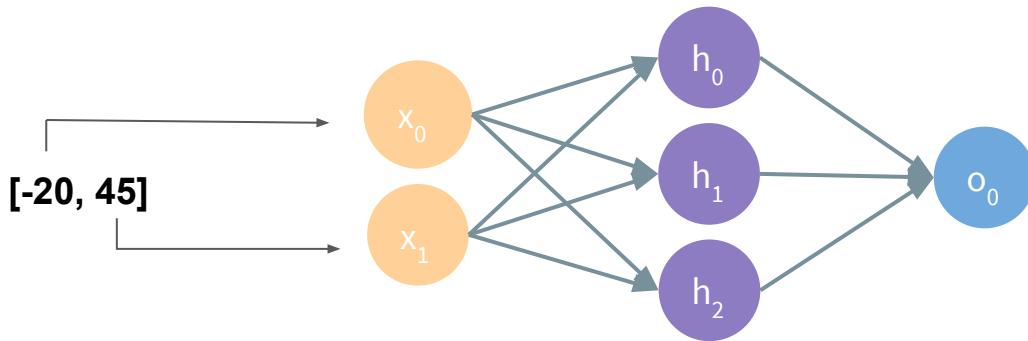
Example Problem: Will my Flight be Delayed?

[-20, 45]

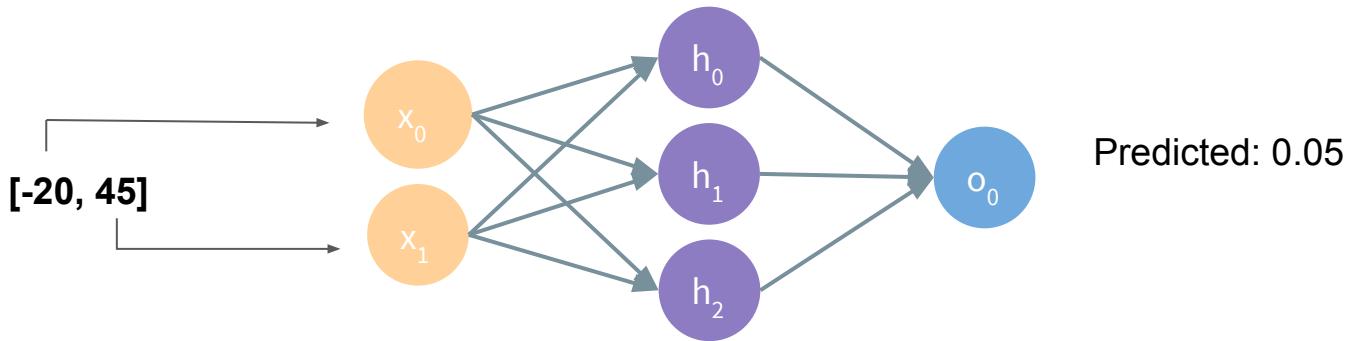


0805	0815	MILANO	•• DELAYED
0845	0855	PARIS	•• DELAYED
0905	0915	NEW YORK	•• DELAYED
0910	0920	FRANKFURT	•• DELAYED
0925	0955	LONDON	•• DELAYED

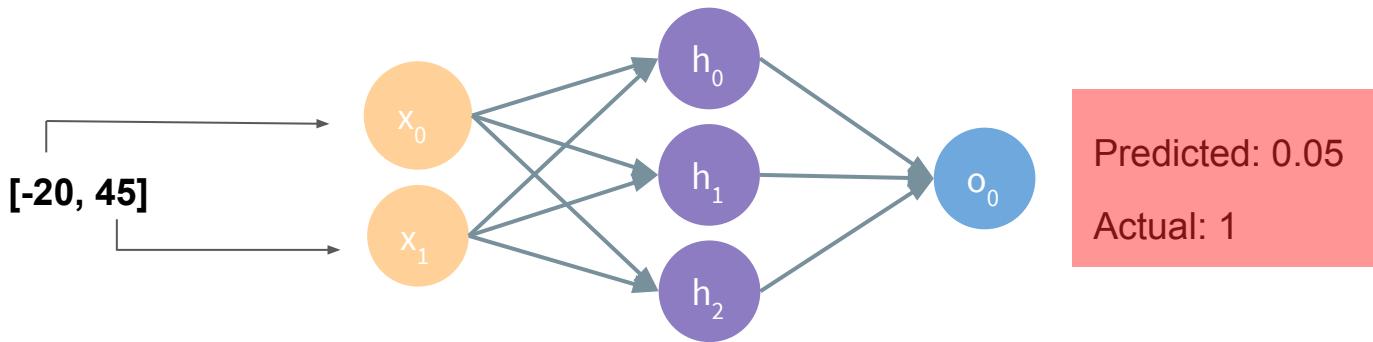
Example Problem: Will my Flight be Delayed?



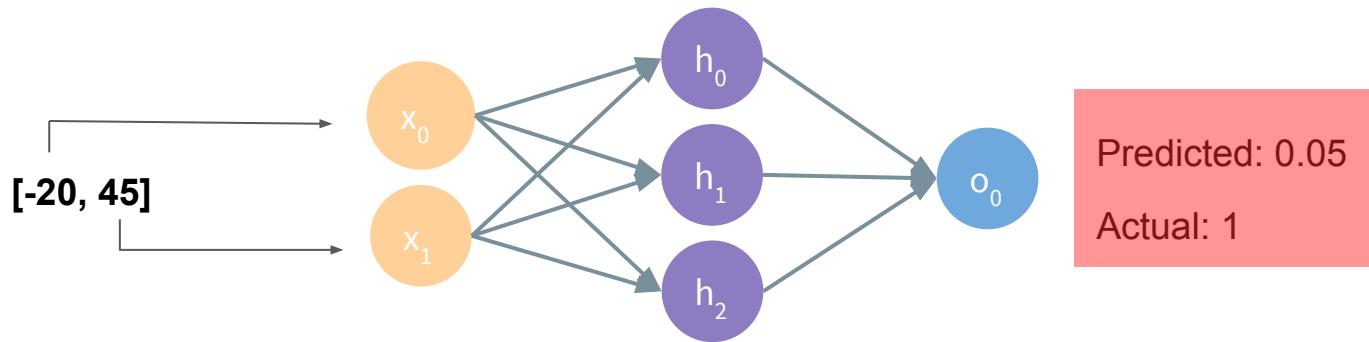
Example Problem: Will my Flight be Delayed?



Example Problem: Will my Flight be Delayed?



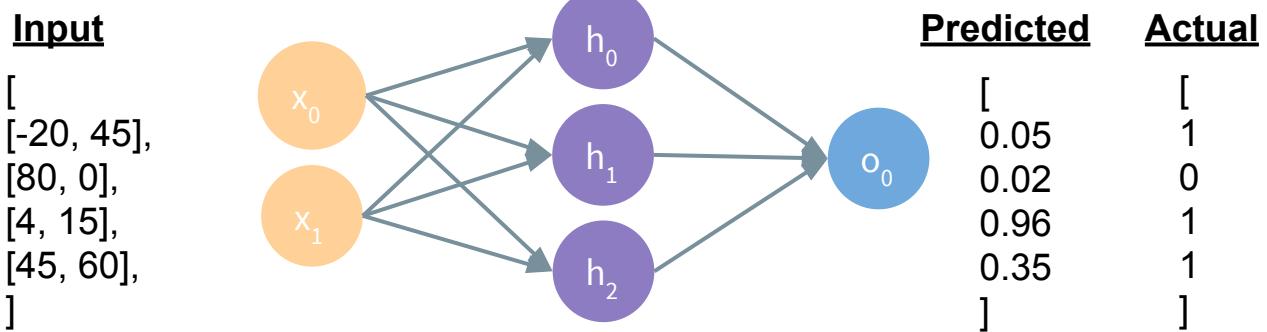
Quantifying Loss



$$loss(f(x^{(i)}; \theta), y^{(i)})$$

Predicted Actual

Total Loss

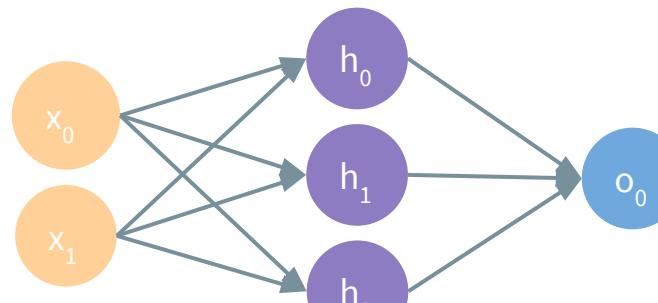


$$\text{total loss} := J(\theta) = \frac{1}{N} \sum_i^N \underline{\text{loss}}(f(x^{(i)}; \theta), \overline{\text{y}}^{(i)})$$

Predicted Actual

Total Loss

Input
[
[-20, 45],
[80, 0],
[4, 15],
[45, 60],
]

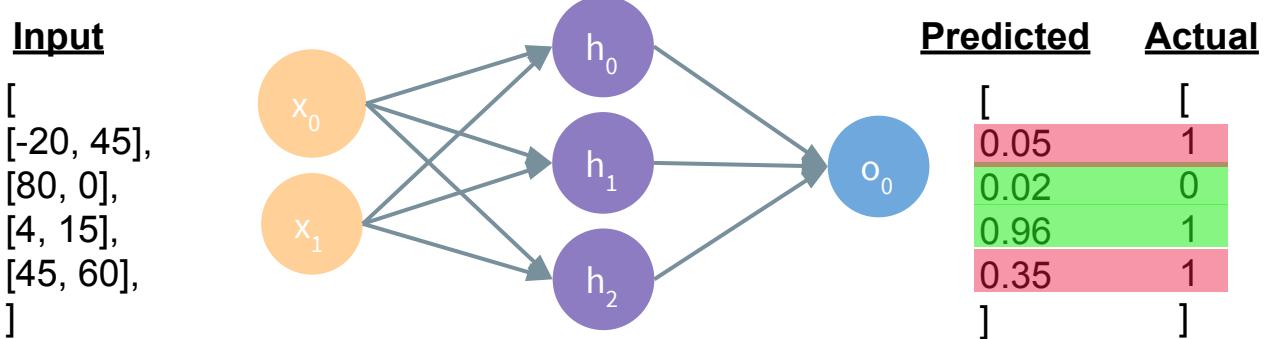


Predicted Actual
[
0.05 1
0.02 0
0.96 1
0.35 1
]
]

$$\text{total loss} := J(\theta) = \frac{1}{N} \sum_i^N \underline{\text{loss}}(f(x^{(i)}; \theta), \overline{\text{y}}^{(i)})$$

Predicted Actual

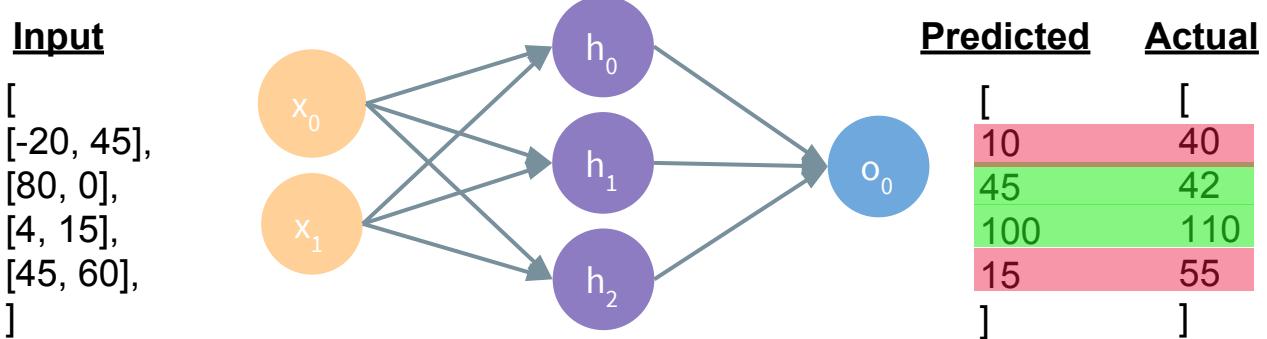
Binary Cross Entropy Loss



$$\text{cross_entropy}(\theta) = \frac{1}{N} \sum_i^N \underline{y^{(i)} \log(f(x^{(i)}; \theta)) + (1 - y^{(i)}) \log(1 - f(x^{(i)}; \theta))}$$

Actual Predicted Actual Predicted

Mean Squared Error (MSE) Loss



$$\text{MSE}(\theta) = \frac{1}{N} \sum_i^N \underline{(f(x^{(i)}; \theta) - y^{(i)})^2}$$

Predicted Actual

Training Neural Networks

Training Neural Networks: Objective

$$\arg_{\theta} \min \frac{1}{N} \sum_i^N \text{loss}(f(x^{(i)}; \theta), y^{(i)})$$

Training Neural Networks: Objective

$$\arg_{\theta} \min \frac{1}{N} \sum_i^N \text{loss}(f(x^{(i)}; \theta), y^{(i)})$$



$$J(\theta)$$



loss function

Training Neural Networks: Objective

$$\arg_{\theta} \min \frac{1}{N} \sum_i^N \text{loss}(f(x^{(i)}; \theta), y^{(i)})$$

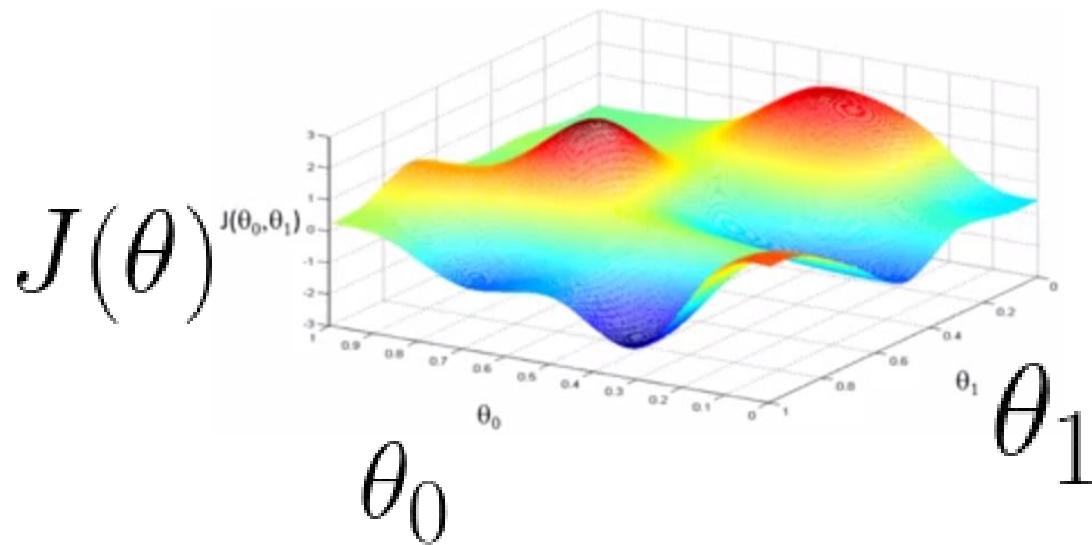


$$J(\theta)$$

$$\theta = W_1, W_2 \dots W_n$$

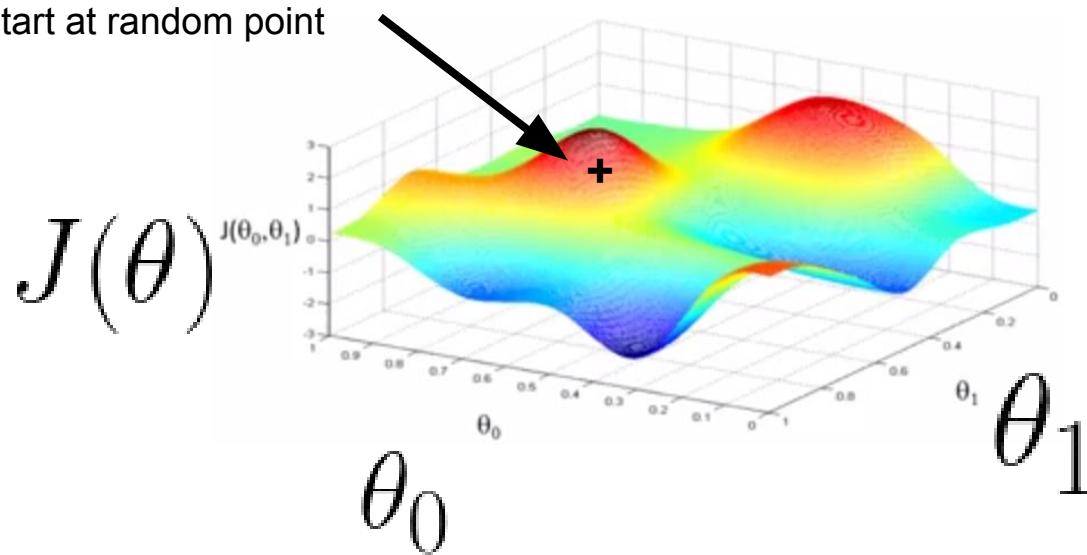


Loss is a **function** of the model's parameters



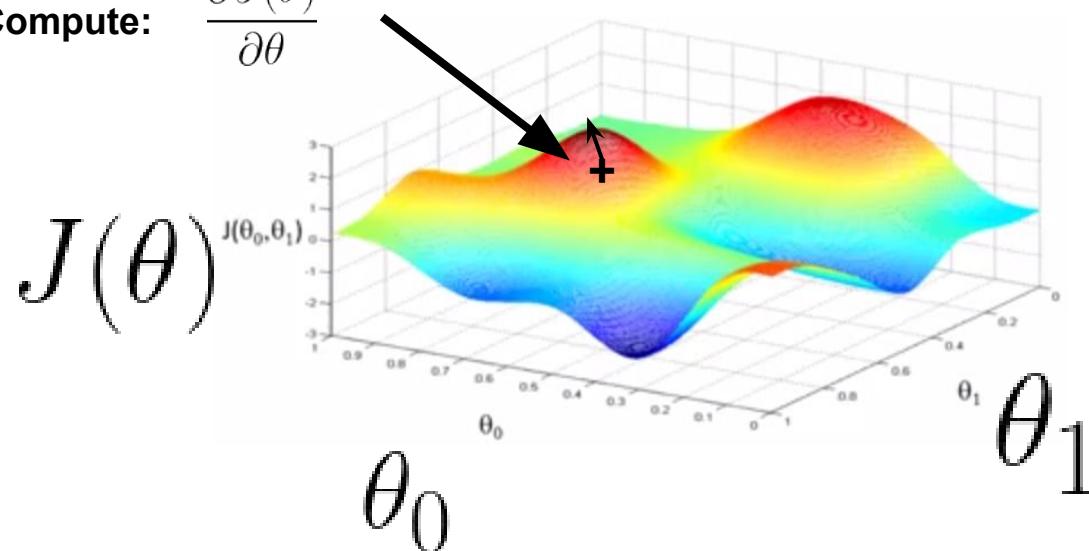
How to minimize loss?

Start at random point



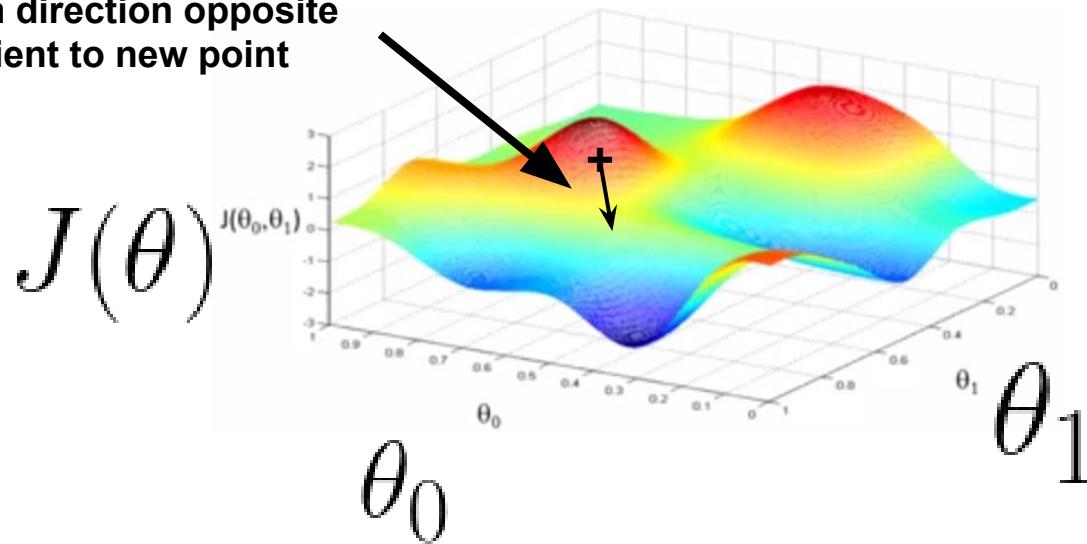
How to minimize loss?

Compute: $\frac{\partial J(\theta)}{\partial \theta}$



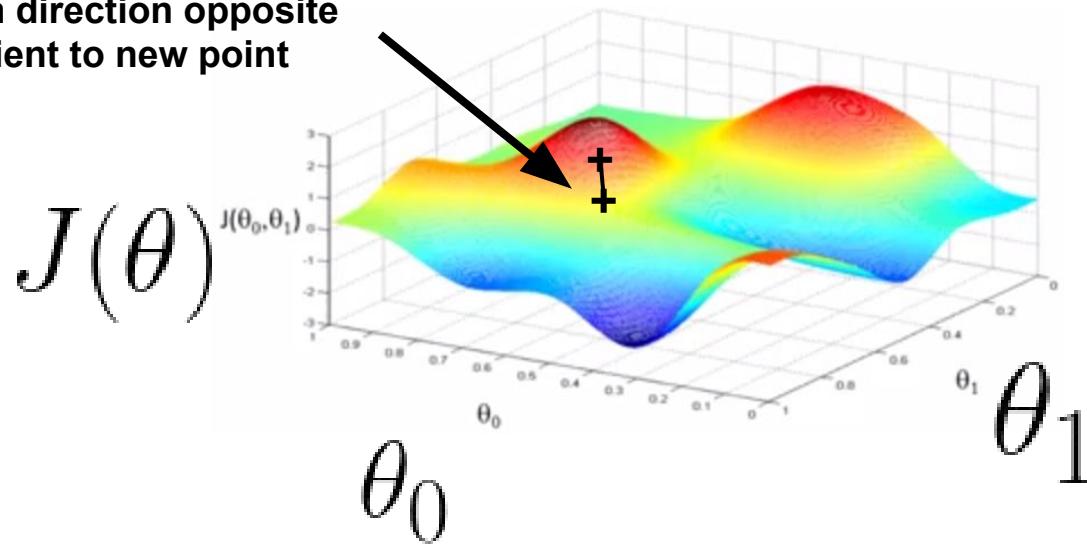
How to minimize loss?

Move in direction opposite
of gradient to new point



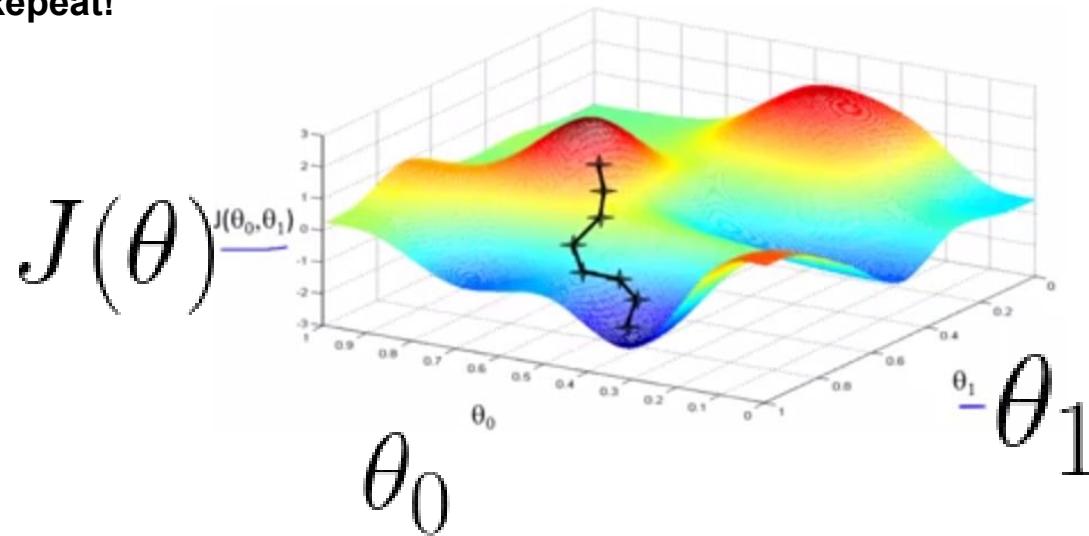
How to minimize loss?

Move in direction opposite
of gradient to new point



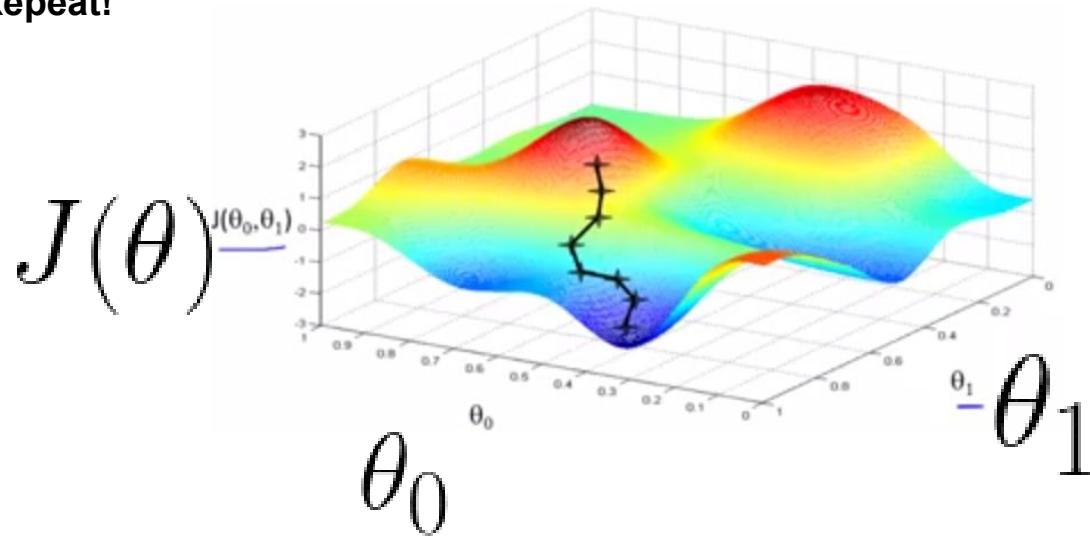
How to minimize loss?

Repeat!



This is called Stochastic Gradient Descent (SGD)

Repeat!



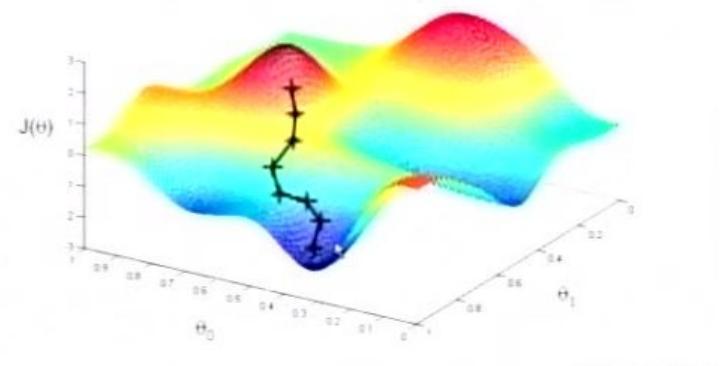
Stochastic Gradient Descent (SGD)

- Initialize θ randomly
- For N Epochs
 - For each training example (x, y) :

- Compute Loss Gradient: $\frac{\partial J(\theta)}{\partial \theta}$

- Update θ with update rule:

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

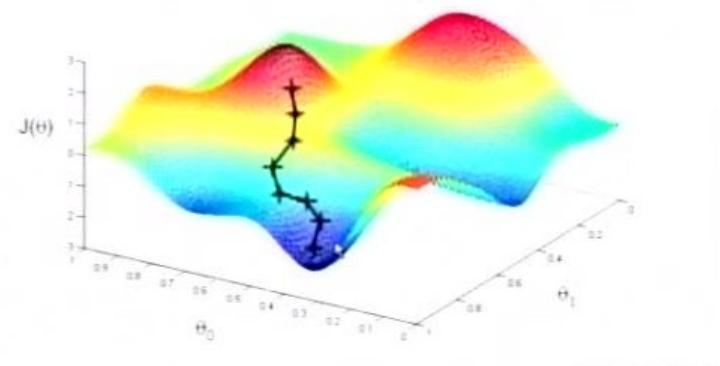


Stochastic Gradient Descent (SGD)

- Initialize θ randomly
- For N Epochs
 - For each training example (x, y) :

- Compute Loss Gradient: $\frac{\partial J(\theta)}{\partial \theta}$
- Update θ with update rule:

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

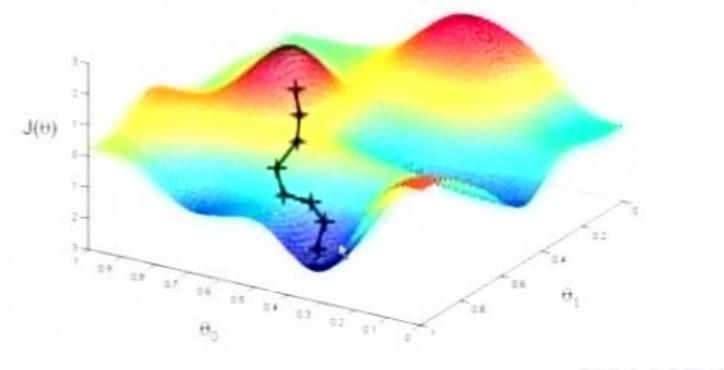


Stochastic Gradient Descent (SGD)

- Initialize θ randomly
- For N Epochs
 - For each training example (x, y) :

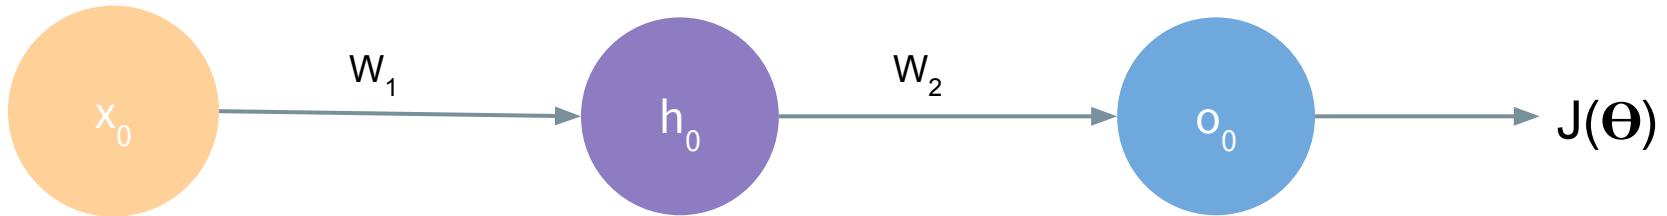
- Compute Loss Gradient: $\frac{\partial J(\theta)}{\partial \theta}$
- Update θ with update rule:

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

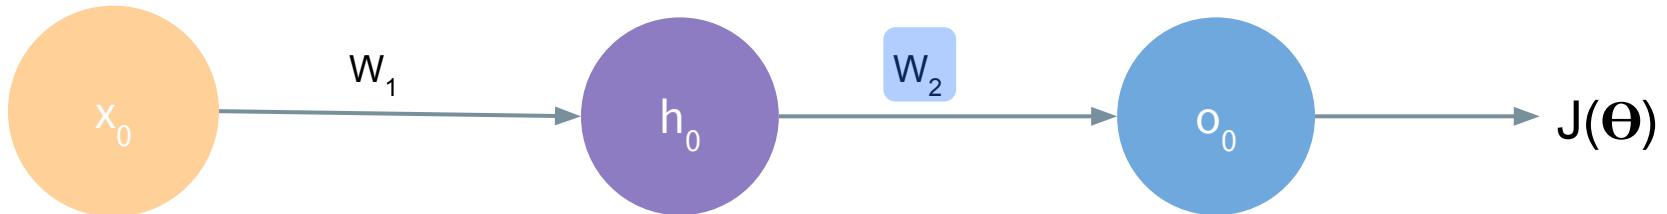


- How to Compute Gradient?

Calculating the Gradient: Backpropagation

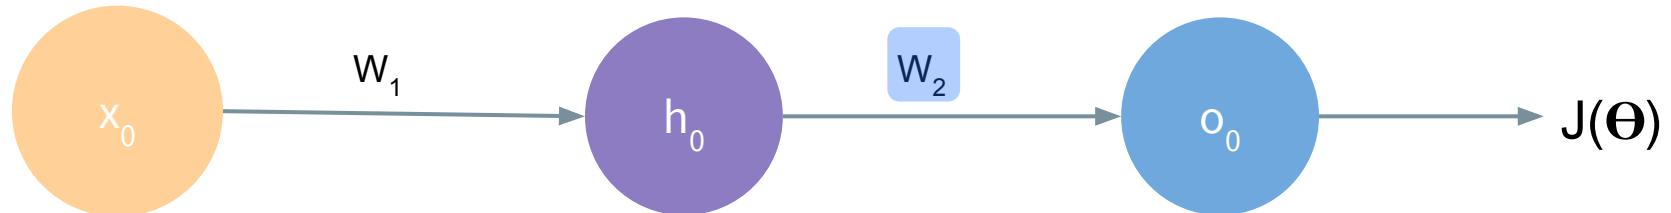


Calculating the Gradient: Backpropagation



$$\frac{\partial J(\theta)}{\partial W_2} =$$

Calculating the Gradient: Backpropagation

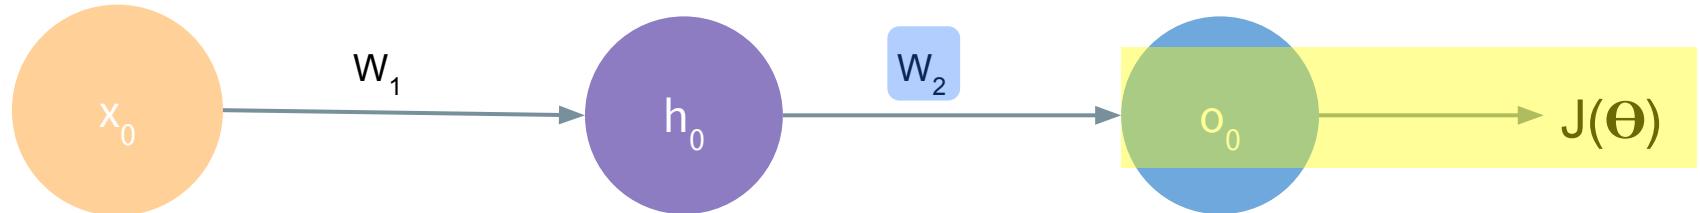


Apply the chain rule

$$\frac{\partial J(\theta)}{\partial W_2} =$$



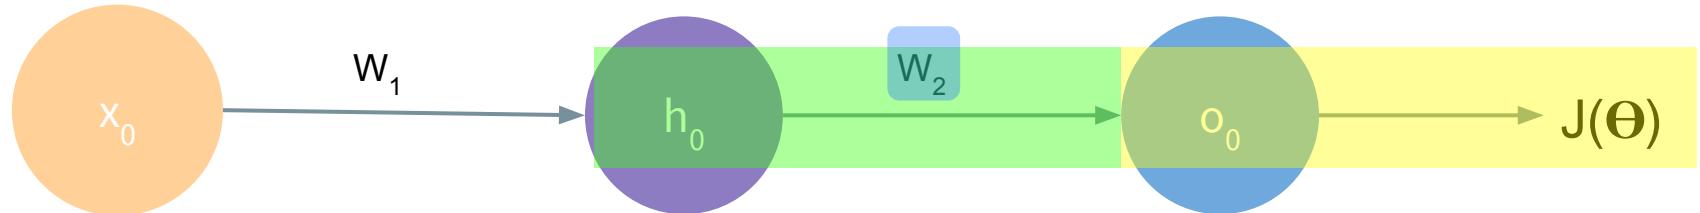
Calculating the Gradient: Backpropagation



Apply the chain rule

$$\frac{\partial J(\theta)}{\partial W_2} = \frac{\partial J(\theta)}{\partial o_0}$$

Calculating the Gradient: Backpropagation



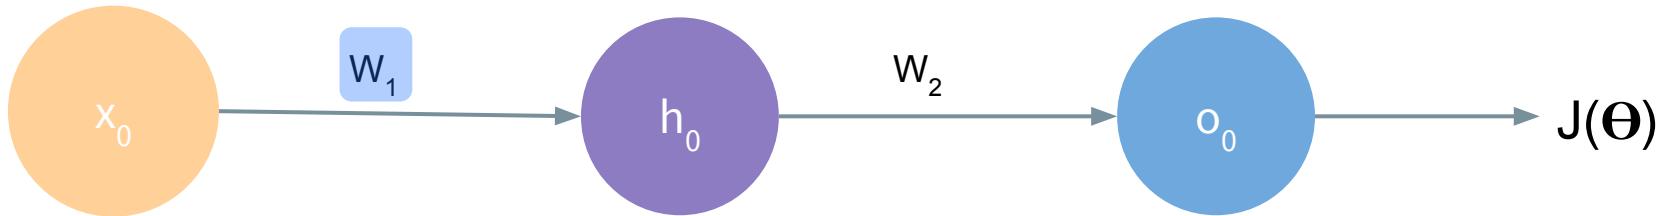
Apply the chain rule

$$\frac{\partial J(\theta)}{\partial W_2} = \frac{\partial J(\theta)}{\partial o_0} * \frac{\partial o_0}{\partial W_2}$$



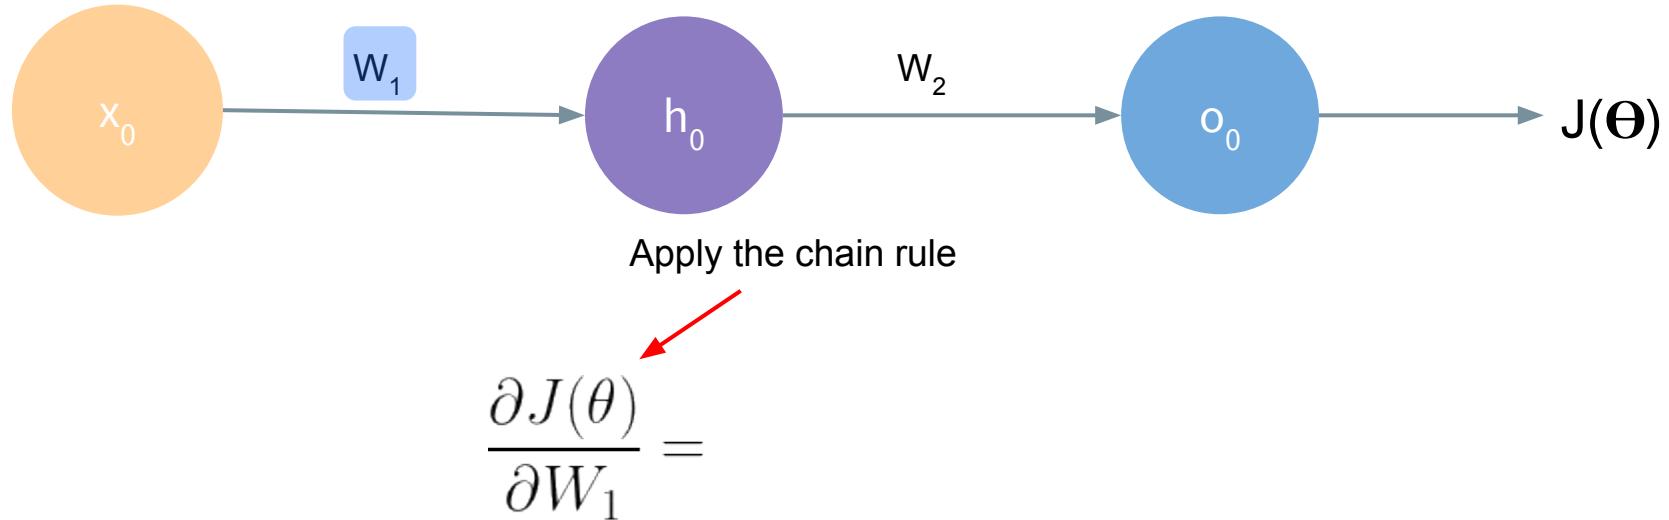
MIT 6.S191 | Intro to Deep Learning | IAP 2017

Calculating the Gradient: Backpropagation

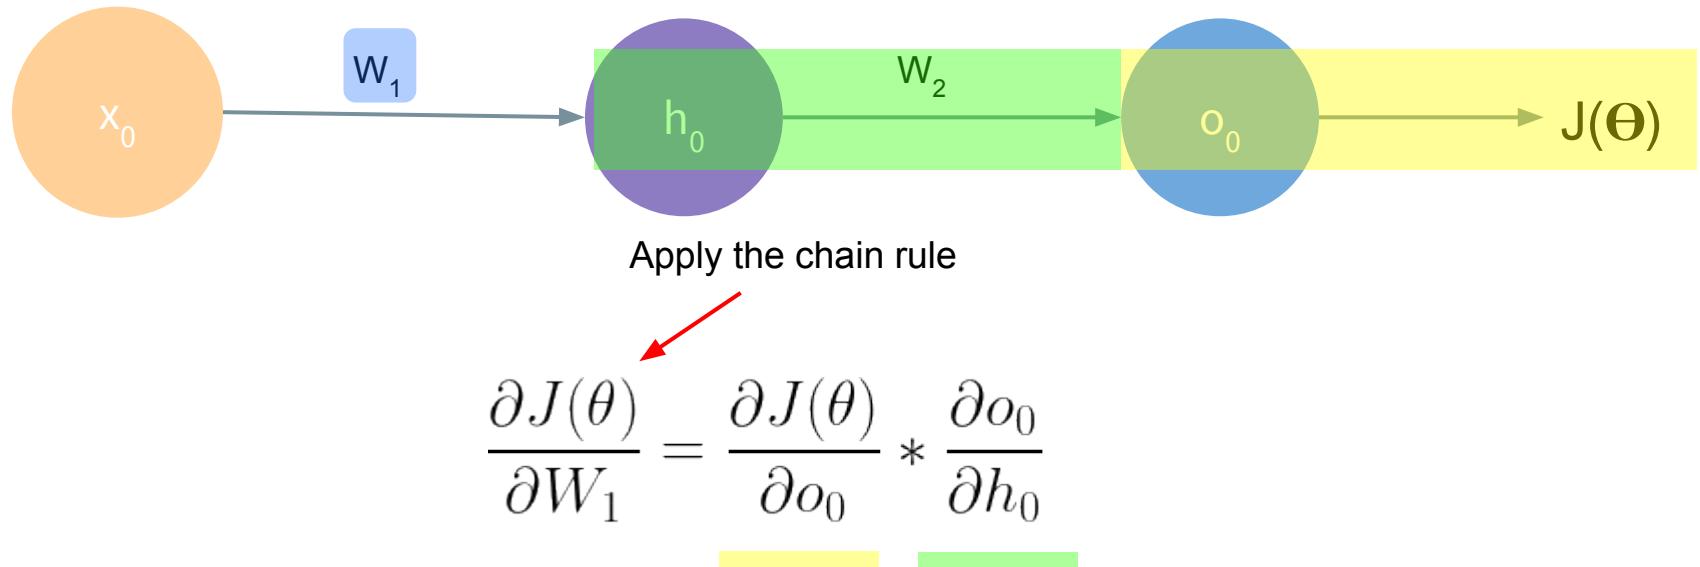


$$\frac{\partial J(\theta)}{\partial W_1} =$$

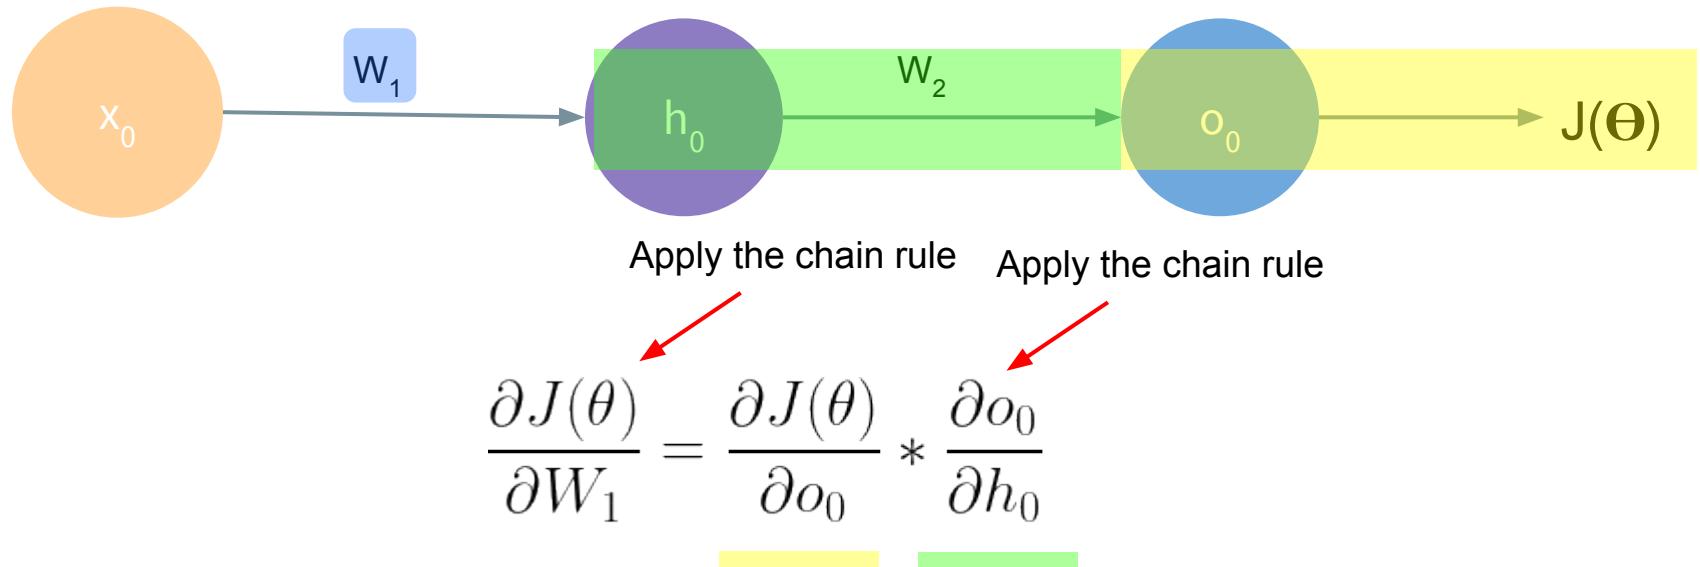
Calculating the Gradient: Backpropagation



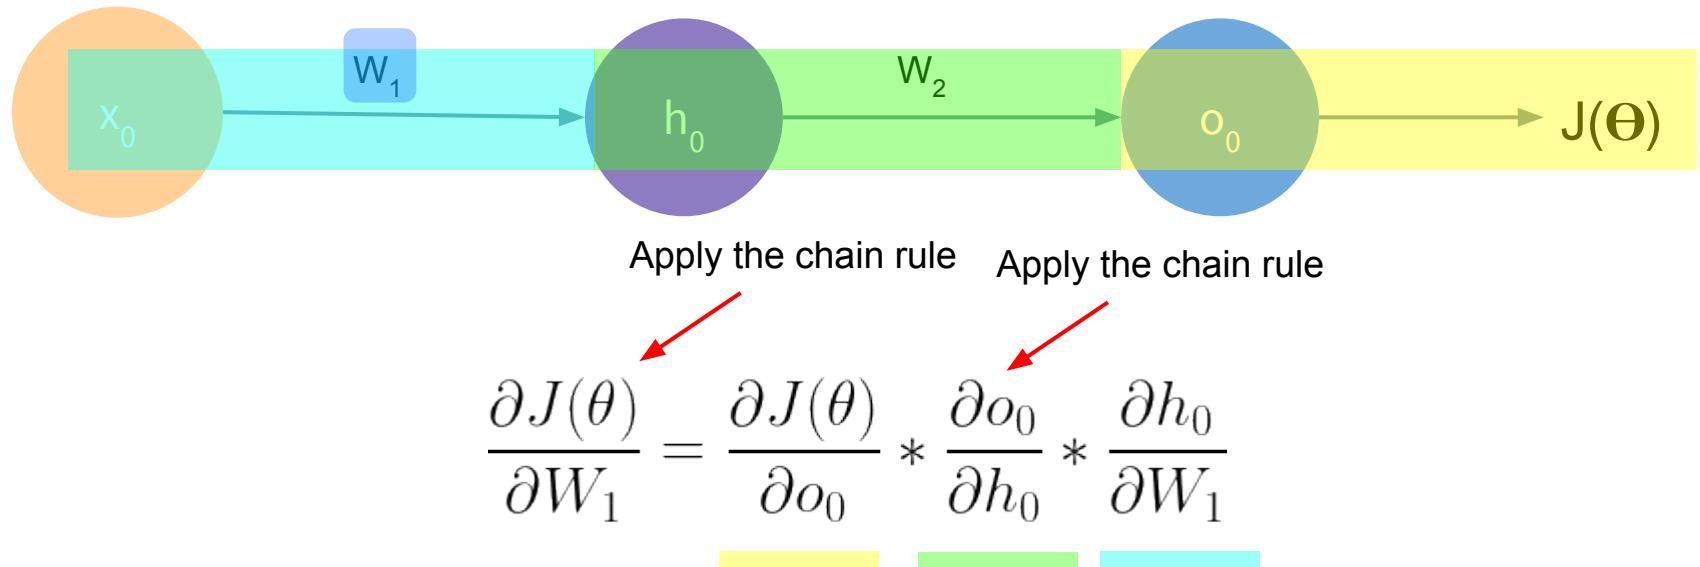
Calculating the Gradient: Backpropagation



Calculating the Gradient: Backpropagation

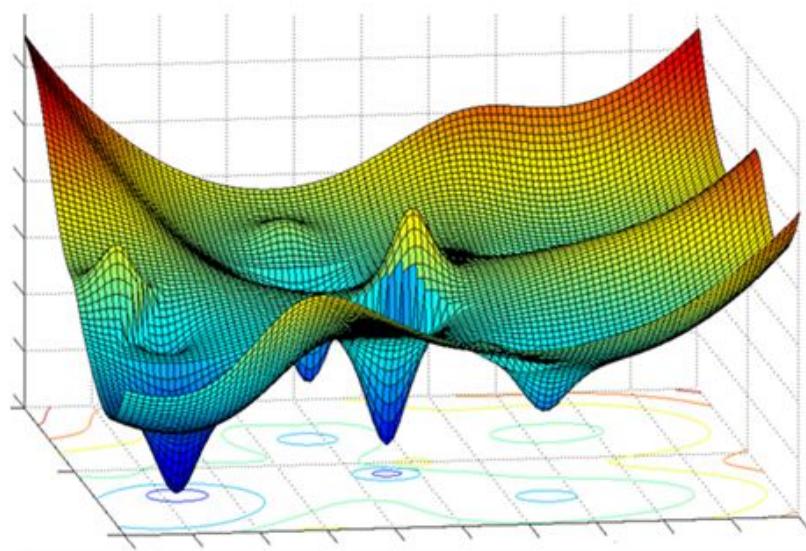


Calculating the Gradient: Backpropagation



Training Neural Networks In Practice

Loss function can be difficult to optimize



Loss function can be difficult to optimize

Update Rule: $\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$

Loss function can be difficult to optimize

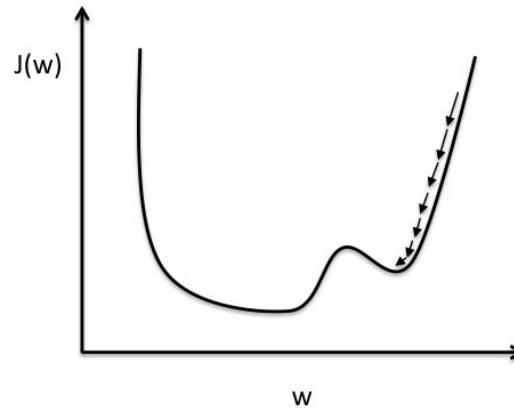
How to Choose Learning Rate?

Update Rule: $\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$



Learning Rate & Optimization

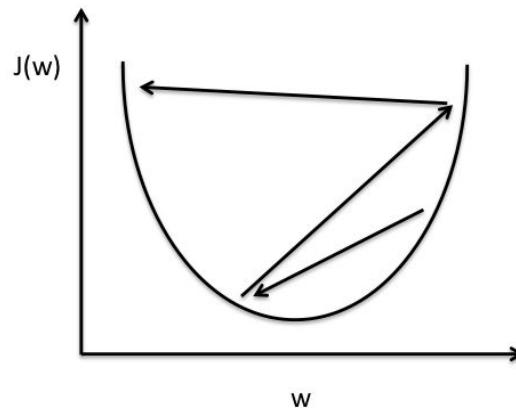
- Small Learning Rate



Small learning rate: Many iterations until convergence and trapping in local minima.

Learning Rate & Optimization

- Large learning rate



Large learning rate: Overshooting.

How to deal with this?

1. Try lots of different learning rates to see what is 'just right'

How to deal with this?

1. Try lots of different learning rates to see what is 'just right'
2. Do something smarter

How to deal with this?

1. Try lots of different learning rates to see what is 'just right'
2. **Do something smarter : Adaptive Learning Rate**

Adaptive Learning Rate

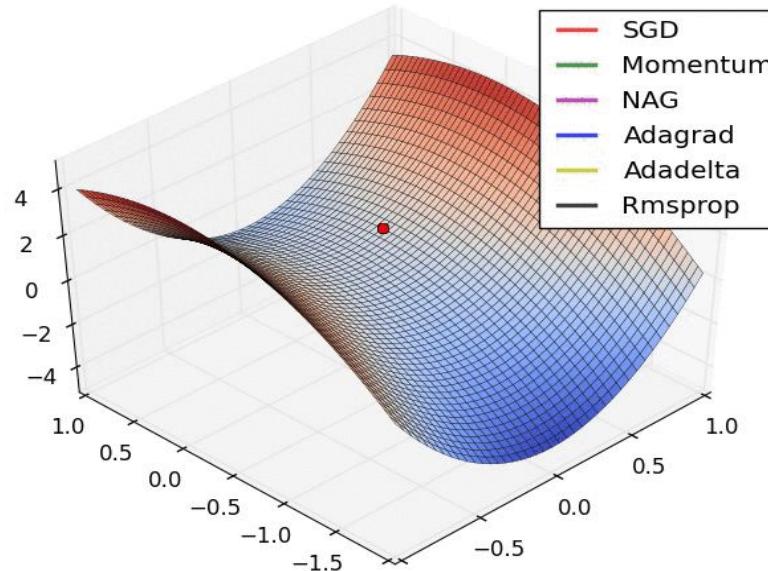
- Learning rate is no longer fixed
- Can be made larger or smaller depending on:
 - how large gradient is
 - how fast learning is happening
 - size of particular weights
 - etc

Adaptive Learning Rate Algorithms

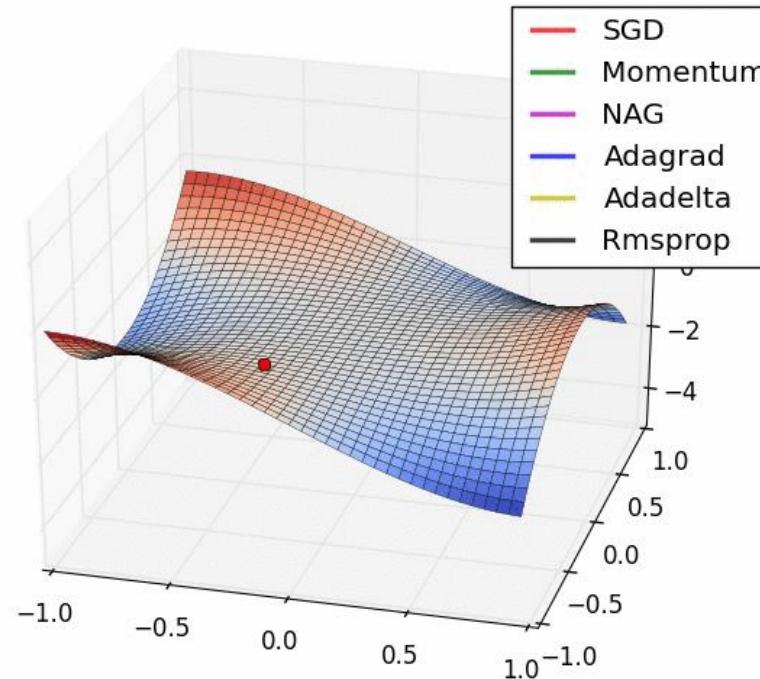
- ADAM
- Momentum
- NAG
- Adagrad
- Adadelta
- RMSProp

For details: check out <http://sebastianruder.com/optimizing-gradient-descent/>

Escaping Saddle Points



Escaping Saddle Points



Training Neural Networks In Practice 2: MiniBatches

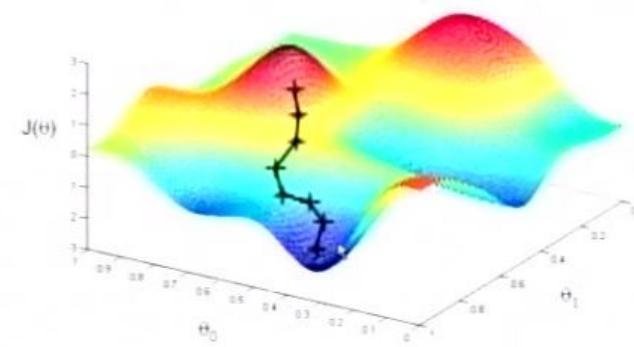
Why is it Stochastic Gradient Descent?

- Initialize θ randomly
- For N Epochs
 - For each training example (x, y) :

- Compute Loss Gradient:
- Update θ with update rule:

$$\frac{\partial J(\theta)}{\partial \theta}$$

Only an estimate of
true gradient!



$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

Minibatches Reduce Gradient Variance

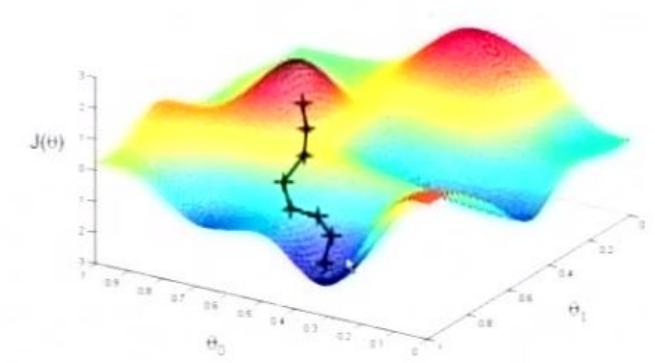
- Initialize θ randomly
- For N Epochs

- For each training **batch** $\{(x_0, y_0), \dots, (x_B, y_B)\}$:

- Compute Loss Gradient:
$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{B} \sum_i^B \frac{\partial J_i(\theta)}{\partial \theta}$$
 - Update θ with update rule:

$$\theta := \theta - \eta \frac{\partial J(\theta)}{\partial \theta}$$

More accurate estimate!

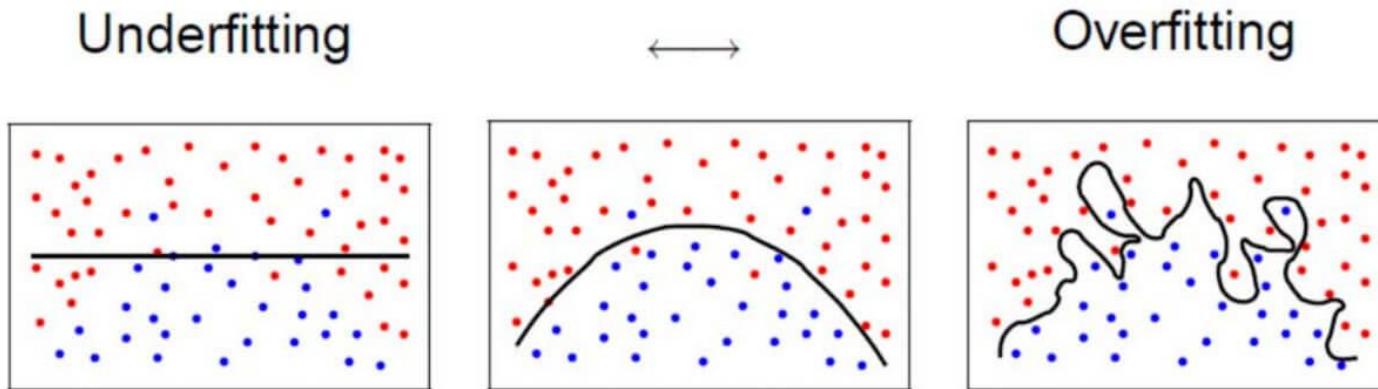


Advantages of Minibatches

- More accurate estimation of gradient
 - Smoother convergence
 - Allows for larger learning rates
- Minibatches lead to fast training!
 - Can parallelize computation + achieve significant speed increases on GPU's

Training Neural Networks In Practice 3: Fighting Overfitting

The Problem of Overfitting

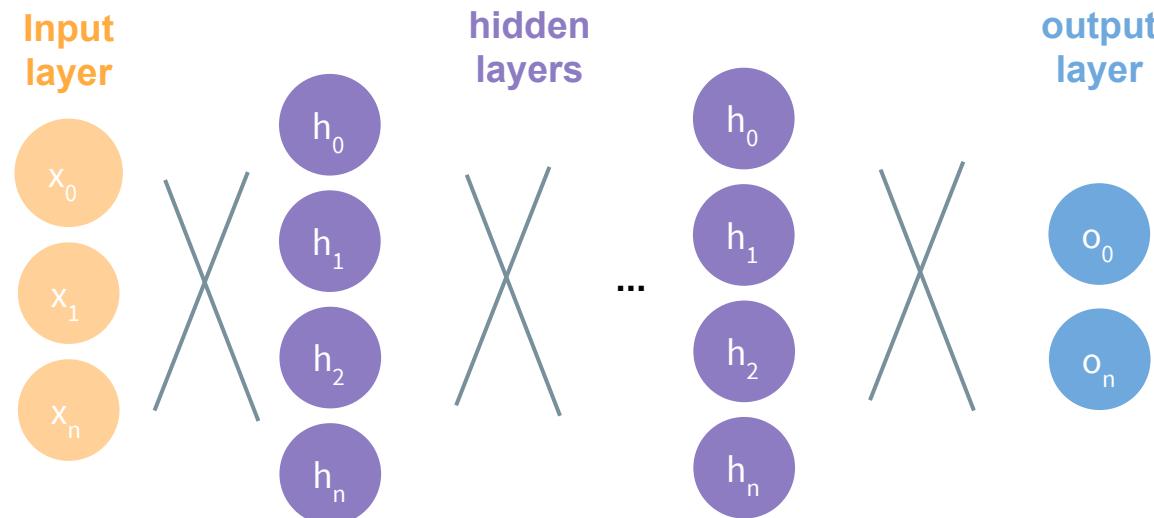


Regularization Techniques

1. Dropout
2. Early Stopping
3. Weight Regularization
4. ...many more

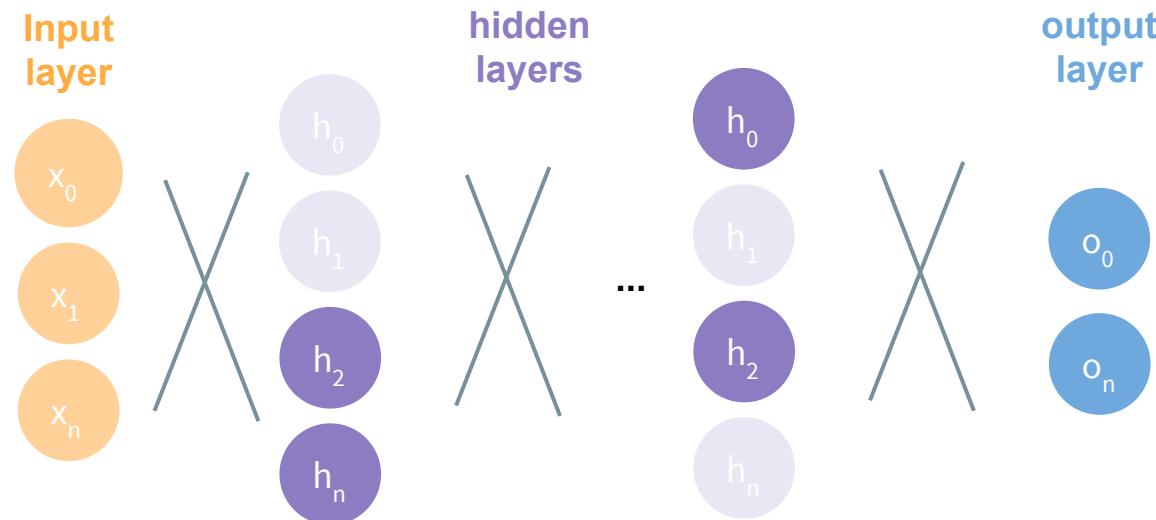
Regularization I: Dropout

- During training, randomly set some activations to 0



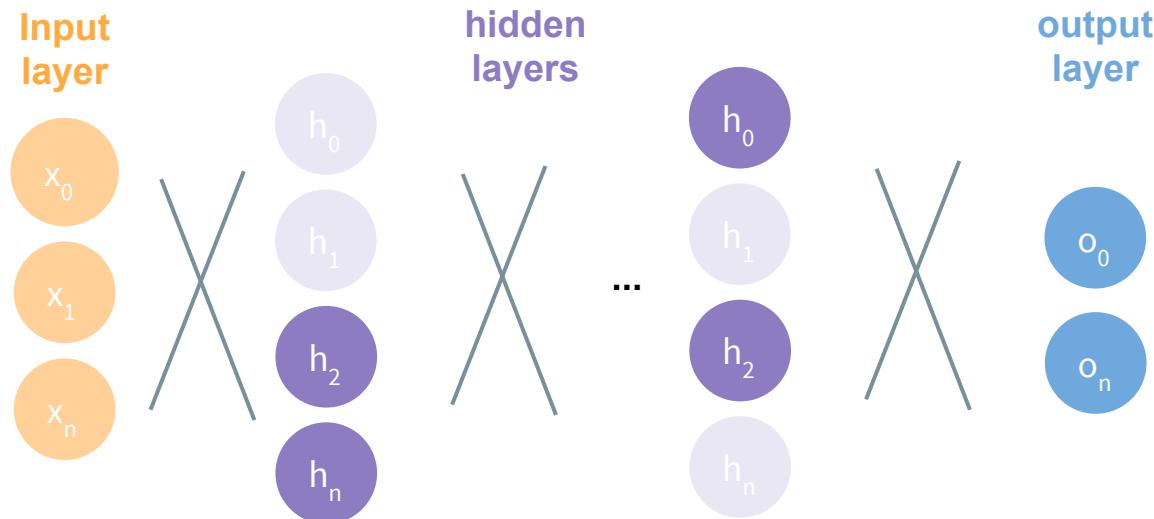
Regularization I: Dropout

- During training, randomly set some activations to 0



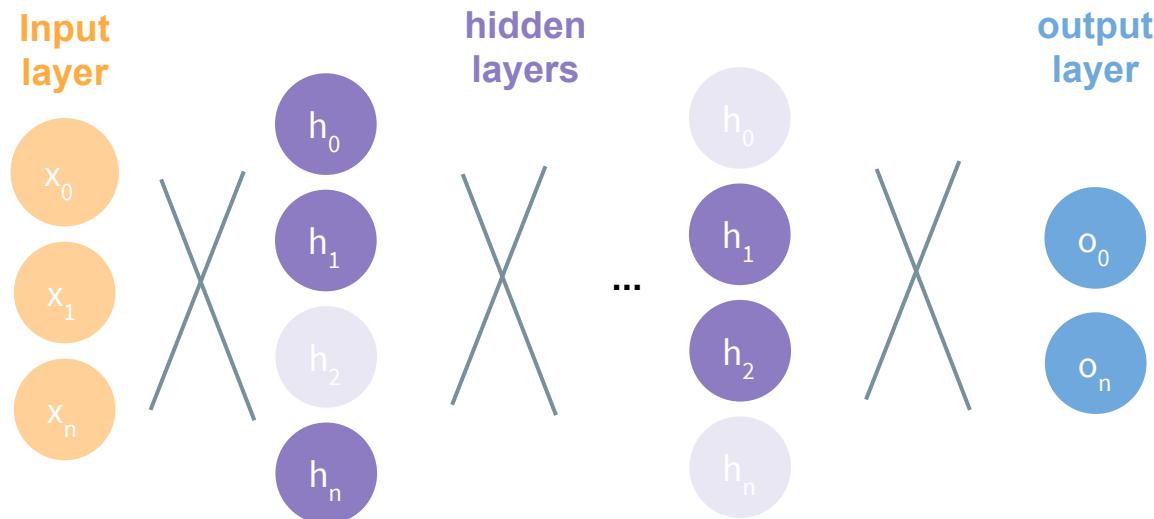
Regularization I: Dropout

- During training, randomly set some activations to 0
 - Typically ‘drop’ 50% of activations in layer
 - Forces network to not rely on any 1 node



Regularization I: Dropout

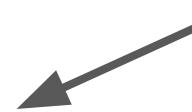
- During training, randomly set some activations to 0
 - Typically ‘drop’ 50% of activations in layer
 - Forces network to not rely on any 1 node



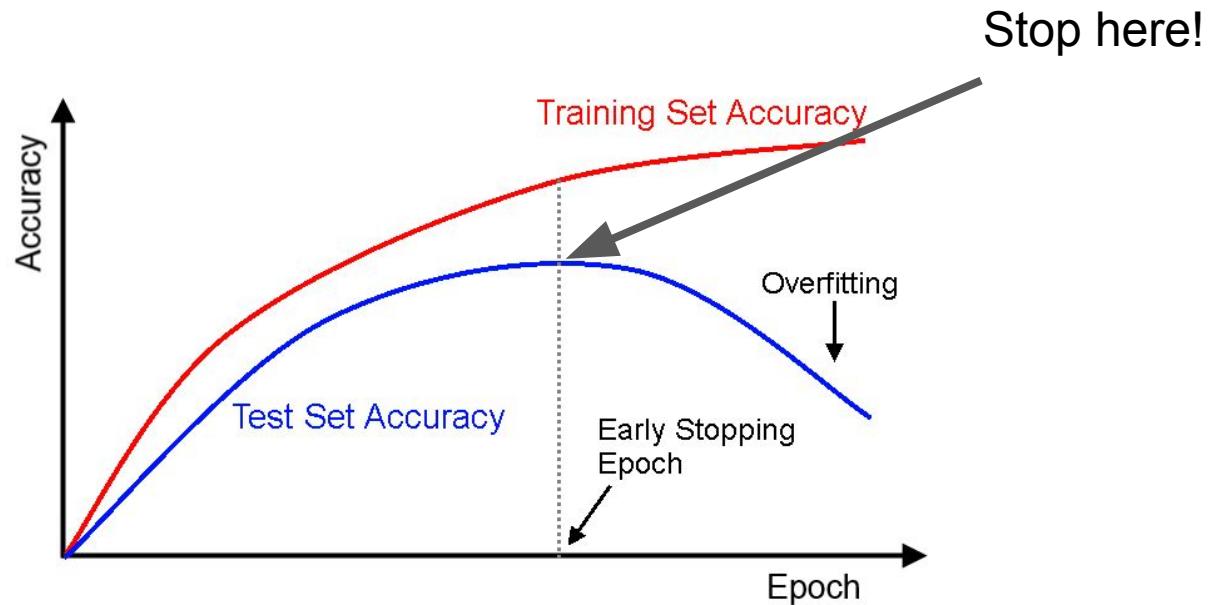
Regularization II: Early Stopping

- Don't give the network time to overfit
- ...
- **Epoch 15:** Train: 85% Validation: 80%
- **Epoch 16:** Train: 87% Validation: 82%
- **Epoch 17:** Train: 90% Validation: 85%
- **Epoch 18:** Train: 95% Validation: 83%
- **Epoch 19:** Train: 97% Validation: 78%
- **Epoch 20:** Train: 98% Validation: 75%

Regularization II: Early Stopping

- Don't give the network time to overfit
 - ...
 - Epoch 15: Train: 85% Validation: 80%
 - Epoch 16: Train: 87% Validation: 82%
 - Epoch 17: Train: 90% Validation: 85%
 - Epoch 18: Train: 95% Validation: 83%
 - Epoch 19: Train: 97% Validation: 78%
 - Epoch 20: Train: 98% Validation: 75%
- 
- Stop here!

Regularization II: Early Stopping



Regularization III: Weight Regularization

- Large weights typically mean model is overfitting
- Add the size of the weights to our loss function
- Perform well on task + keep weights small

Regularization III: Weight Regularization

- Large weights typically mean model is overfitting
- Add the size of the weights to our loss function
- Perform well on task + keep weights small

$$\arg_{\theta} \min \frac{1}{T} \sum_t \text{loss}(f(x^{(t)}; \theta), y^{(t)})$$

Regularization III: Weight Regularization

- Large weights typically mean model is overfitting
- Add the size of the weights to our loss function
- Perform well on task + keep weights small

$$\arg_{\theta} \min \frac{1}{T} \sum_t \text{loss}(f(x^{(t)}; \theta), y^{(t)}) + \lambda \sum_i (\theta_i)^2$$

Regularization III: Weight Regularization

- Large weights typically mean model is overfitting
- Add the size of the weights to our loss function
- Perform well on task + keep weights small

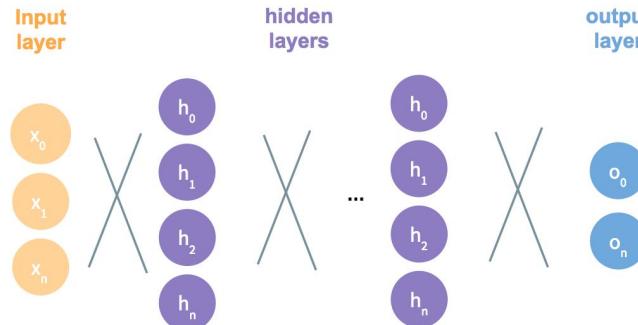
$$\arg_{\theta} \min \frac{1}{T} \sum_t \text{loss}(f(x^{(t)}; \theta), y^{(t)}) + \lambda \sum_i (\theta_i)^2$$



$$J(\theta)$$

Core Fundamentals Review

- Perceptron Classifier
- Stacking Perceptrons to form neural networks
- How to formulate problems with neural networks
- Train neural networks with backpropagation
- Techniques for improving training of deep neural networks



Questions?